

# Package ‘ggmlR’

January 28, 2026

**Type** Package

**Title** GGML Tensor Operations for Machine Learning

**Version** 0.4.1

**Description** An R port of the 'GGML' library providing efficient tensor operations for machine learning. Implements core functionality including element-wise arithmetic, tensor manipulation, and multi-backend support. Brings lightweight, performance-oriented design to 'R' for fast numerical computation and model prototyping. See <<https://github.com/ggml-org/ggml>> for more information about the underlying library.

**License** MIT + file LICENSE

**URL** <https://github.com/Zabis13/ggmlR>

**BugReports** <https://github.com/Zabis13/ggmlR/issues>

**Encoding** UTF-8

**SystemRequirements** C++17, GNU make

**Suggests** testthat (>= 3.0.0)

**RoxygenNote** 7.3.3

**Config/testthat.edition** 3

**NeedsCompilation** yes

**Author** Yuri Baramykov [aut, cre]

**Maintainer** Yuri Baramykov <lbsbmsu@mail.ru>

## R topics documented:

ggmlR-package . . . . .	6
ggml_abs . . . . .	7
ggml_add . . . . .	7
ggml_add1 . . . . .	8
ggml_are_same_shape . . . . .	9
ggml_argmax . . . . .	9

ggml_argsort . . . . .	10
ggml_backend_alloc_ctx_tensors . . . . .	11
ggml_backend_buffer_free . . . . .	11
ggml_backend_buffer_get_size . . . . .	12
ggml_backend_buffer_name . . . . .	12
ggml_backend_cpu_init . . . . .	13
ggml_backend_cpu_set_n_threads . . . . .	13
ggml_backend_free . . . . .	14
ggml_backend_graph_compute . . . . .	14
ggml_backend_name . . . . .	15
ggml_backend_sched_alloc_graph . . . . .	15
ggml_backend_sched_free . . . . .	16
ggml_backend_sched_get_backend . . . . .	16
ggml_backend_sched_get_n_backends . . . . .	17
ggml_backend_sched_get_n_copies . . . . .	17
ggml_backend_sched_get_n_splits . . . . .	18
ggml_backend_sched_get_tensor_backend . . . . .	18
ggml_backend_sched_graph_compute . . . . .	19
ggml_backend_sched_graph_compute_async . . . . .	20
ggml_backend_sched_new . . . . .	20
ggml_backend_sched_reserve . . . . .	21
ggml_backend_sched_reset . . . . .	22
ggml_backend_sched_set_tensor_backend . . . . .	23
ggml_backend_sched_synchronize . . . . .	23
ggml_backend_tensor_get_data . . . . .	24
ggml_backend_tensor_set_data . . . . .	24
ggml_build_forward_expand . . . . .	25
ggml_ceil . . . . .	26
ggml_clamp . . . . .	26
ggml_concat . . . . .	27
ggml_cont . . . . .	28
ggml_conv_1d . . . . .	28
ggml_conv_2d . . . . .	29
ggml_conv_transpose_1d . . . . .	29
ggml_cos . . . . .	30
ggml_cpu_add . . . . .	31
ggml_cpu_mul . . . . .	31
ggml_cpy . . . . .	32
ggml_cycles . . . . .	33
ggml_cycles_per_ms . . . . .	34
ggml_diag . . . . .	34
ggml_diag_mask_inf . . . . .	35
ggml_diag_mask_inf_inplace . . . . .	36
ggml_diag_mask_zero . . . . .	37
ggml_div . . . . .	37
ggml_dup . . . . .	38
ggml_dup_tensor . . . . .	39
ggml_element_size . . . . .	39

ggml_elu . . . . .	40
ggml_estimate_memory . . . . .	41
ggml_exp . . . . .	41
ggml_flash_attn_back . . . . .	42
ggml_flash_attn_ext . . . . .	42
ggml_floor . . . . .	44
ggml_free . . . . .	44
ggml_gallocr_alloc_graph . . . . .	45
ggml_gallocr_free . . . . .	46
ggml_gallocr_get_buffer_size . . . . .	46
ggml_gallocr_new . . . . .	47
ggml_gallocr_reserve . . . . .	47
ggml_geglu . . . . .	48
ggml_geglu_quick . . . . .	49
ggml_geglu_split . . . . .	49
ggml_gelu . . . . .	50
ggml_gelu_erf . . . . .	50
ggml_gelu_quick . . . . .	51
ggml_get_f32 . . . . .	52
ggml_get_i32 . . . . .	53
ggml_get_max_tensor_size . . . . .	53
ggml_get_mem_size . . . . .	54
ggml_get_name . . . . .	55
ggml_get_no_alloc . . . . .	55
ggml_get_n_threads . . . . .	56
ggml_get_rows . . . . .	56
ggml_get_rows_back . . . . .	57
ggml_glu . . . . .	58
GGML_GLU_OP_REGLU . . . . .	59
ggml_glu_split . . . . .	60
ggml_graph_compute . . . . .	60
ggml_graph_compute_with_ctx . . . . .	61
ggml_graph_dump_dot . . . . .	62
ggml_graph_get_tensor . . . . .	63
ggml_graph_node . . . . .	63
ggml_graph_n_nodes . . . . .	64
ggml_graph_overhead . . . . .	64
ggml_graph_print . . . . .	65
ggml_graph_reset . . . . .	65
ggml_group_norm . . . . .	66
ggml_group_norm_inplace . . . . .	66
ggml_hardsigmoid . . . . .	67
ggml_hardswish . . . . .	67
ggml_im2col . . . . .	68
ggml_init . . . . .	69
ggml_init_auto . . . . .	70
ggml_is_available . . . . .	70
ggml_is_contiguous . . . . .	71

ggml_is_permuted . . . . .	71
ggml_is_transposed . . . . .	72
ggml_l2_norm . . . . .	73
ggml_l2_norm_inplace . . . . .	73
ggml_leaky_relu . . . . .	74
ggml_log . . . . .	74
ggml_mean . . . . .	75
ggml_mul . . . . .	75
ggml_mul_mat . . . . .	76
ggml_mul_mat_id . . . . .	77
ggml nbytes . . . . .	78
ggml_neg . . . . .	79
ggml_nelements . . . . .	79
ggml_new_f32 . . . . .	80
ggml_new_i32 . . . . .	81
ggml_new_tensor . . . . .	82
ggml_new_tensor_1d . . . . .	82
ggml_new_tensor_2d . . . . .	83
ggml_new_tensor_3d . . . . .	84
ggml_new_tensor_4d . . . . .	85
ggml_norm . . . . .	85
ggml_norm_inplace . . . . .	86
ggml_nrows . . . . .	86
ggml_n_dims . . . . .	87
ggml_out_prod . . . . .	87
ggml_pad . . . . .	88
ggml_permute . . . . .	89
ggml_pool_1d . . . . .	90
ggml_pool_2d . . . . .	91
ggml_print_mem_status . . . . .	91
ggml_print_objects . . . . .	92
ggml_quantize_chunk . . . . .	93
ggml_quantize_free . . . . .	93
ggml_quantize_init . . . . .	94
ggml_quantize_requires_imatrix . . . . .	94
ggml_reglu . . . . .	95
ggml_reglu_split . . . . .	96
ggml_relu . . . . .	96
ggml_repeat . . . . .	97
ggml_repeat_back . . . . .	97
ggml_reset . . . . .	98
ggml_reshape_1d . . . . .	99
ggml_reshape_2d . . . . .	99
ggml_reshape_3d . . . . .	100
ggml_reshape_4d . . . . .	100
ggml_rms_norm . . . . .	101
ggml_rms_norm_back . . . . .	101
ggml_rms_norm_inplace . . . . .	102

ggml_rope . . . . .	102
ggml_rope_ext . . . . .	103
ggml_rope_ext_back . . . . .	105
ggml_rope_inplace . . . . .	106
ggml_round . . . . .	107
ggml_scale . . . . .	107
ggml_set . . . . .	108
ggml_set_1d . . . . .	108
ggml_set_2d . . . . .	109
ggml_set_f32 . . . . .	109
ggml_set_i32 . . . . .	110
ggml_set_name . . . . .	111
ggml_set_no_alloc . . . . .	111
ggml_set_n_threads . . . . .	112
ggml_set_zero . . . . .	113
ggml_sgn . . . . .	113
ggml_sigmoid . . . . .	114
ggml_silu . . . . .	115
ggml_silu_back . . . . .	115
ggml_sin . . . . .	116
ggml_softplus . . . . .	116
ggml_soft_max . . . . .	117
ggml_soft_max_ext . . . . .	118
ggml_soft_max_ext_back . . . . .	119
ggml_soft_max_inplace . . . . .	119
GGML_SORT_ORDER_ASC . . . . .	120
ggml_sqrt . . . . .	120
ggml_sqrt . . . . .	121
ggml_step . . . . .	121
ggml_sub . . . . .	122
ggml_sum . . . . .	123
ggml_sum_rows . . . . .	123
ggml_swiglu . . . . .	124
ggml_swiglu_split . . . . .	125
ggml_tanh . . . . .	125
ggml_tensor_overhead . . . . .	126
ggml_tensor_shape . . . . .	126
ggml_tensor_type . . . . .	127
ggml_test . . . . .	128
ggml_time_init . . . . .	128
ggml_time_ms . . . . .	129
ggml_time_us . . . . .	129
ggml_top_k . . . . .	130
ggml_transpose . . . . .	131
GGML_TYPE_F32 . . . . .	131
ggml_type_size . . . . .	132
ggml_upscale . . . . .	133
ggml_used_mem . . . . .	134

ggml_version . . . . .	134
ggml_view_1d . . . . .	135
ggml_view_2d . . . . .	136
ggml_view_3d . . . . .	136
ggml_view_4d . . . . .	137
ggml_view_tensor . . . . .	138
ggml_vulkan_available . . . . .	138
ggml_vulkan_backend_name . . . . .	139
ggml_vulkan_device_count . . . . .	139
ggml_vulkan_device_description . . . . .	140
ggml_vulkan_device_memory . . . . .	141
ggml_vulkan_free . . . . .	141
ggml_vulkan_init . . . . .	142
ggml_vulkan_is_backend . . . . .	143
ggml_vulkan_list_devices . . . . .	143
ggml_vulkan_status . . . . .	144
ggml_with_temp_ctx . . . . .	144
rope_types . . . . .	145

## Description

An R port of the 'GGML' library providing efficient tensor operations for machine learning. Implements core functionality including element-wise arithmetic, tensor manipulation, and multi-backend support. Brings lightweight, performance-oriented design to 'R' for fast numerical computation and model prototyping. See <https://github.com/ggml-org/ggml> for more information about the underlying library.

## Author(s)

**Maintainer:** Yuri Baramykov <1bsbmsu@mail.ru>

## See Also

Useful links:

- <https://github.com/Zabis13/ggmlR>
- Report bugs at <https://github.com/Zabis13/ggmlR/issues>

---

ggml_abs	<i>Absolute Value (Graph)</i>
----------	-------------------------------

---

**Description**

Creates a graph node for element-wise absolute value:  $|x|$

**Usage**

```
ggml_abs(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the abs operation

---

ggml_add	<i>Add tensors</i>
----------	--------------------

---

**Description**

Creates a graph node for element-wise addition. Must be computed using ggml\_build\_forward\_expand() and ggml\_graph\_compute().

**Usage**

```
ggml_add(ctx, a, b)
```

```
ggml_add(ctx, a, b)
```

**Arguments**

ctx	GGML context
a	First tensor
b	Second tensor (same shape as a)

**Value**

Tensor representing the addition operation

Tensor representing the addition operation

## Examples

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
ggml_set_f32(b, c(5, 4, 3, 2, 1))
result <- ggml_add(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
ggml_set_f32(b, c(5, 4, 3, 2, 1))
result <- ggml_add(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

## ggml\_add1

### Add Scalar to Tensor (Graph)

## Description

Creates a graph node for adding a scalar (1-element tensor) to all elements of a tensor. This is more efficient than creating a full tensor of the same value.

## Usage

```
ggml_add1(ctx, a, b)
```

## Arguments

ctx	GGML context
a	Input tensor
b	Scalar tensor (1-element tensor)

## Value

Tensor representing the operation  $a + b$  (broadcasted)

## Examples

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
scalar <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 1)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
ggml_set_f32(scalar, 10)
result <- ggml_add1(ctx, a, scalar)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

ggml\_are\_same\_shape     *Compare Tensor Shapes*

## Description

Checks if two tensors have the same shape.

## Usage

```
ggml_are_same_shape(a, b)
```

## Arguments

a	First tensor
b	Second tensor

## Value

TRUE if shapes are identical, FALSE otherwise

ggml\_argmax     *Argmax (Graph)*

## Description

Creates a graph node that finds the index of the maximum value. CRITICAL for token generation in LLMs.

## Usage

```
ggml_argmax(ctx, a)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor

**Value**

Tensor with argmax indices

`ggml_argsort`

*Argsort - Get Sorting Indices (Graph)*

**Description**

Returns indices that would sort the tensor rows. Each row is sorted independently.

**Usage**

```
ggml_argsort(ctx, a, order = GGML_SORT_ORDER_ASC)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor to sort (F32)
<code>order</code>	Sort order: GGML_SORT_ORDER_ASC (0) or GGML_SORT_ORDER_DESC (1)

**Value**

Tensor of I32 indices that would sort each row

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
# Create tensor with values to sort
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(3, 1, 4, 1, 5))
# Get indices for ascending sort
indices <- ggml_argsort(ctx, a, GGML_SORT_ORDER_ASC)
graph <- ggml_build_forward_expand(ctx, indices)
ggml_graph_compute(ctx, graph)
result <- ggml_get_i32(indices)
# result: [1, 3, 0, 2, 4] (0-indexed positions for sorted order)
ggml_free(ctx)

## End(Not run)
```

---

ggml\_backend\_alloc\_ctx\_tensors  
*Allocate Context Tensors to Backend*

---

### Description

Allocates all tensors in a GGML context to a specific backend. Returns a buffer that must be freed when no longer needed.

### Usage

```
ggml_backend_alloc_ctx_tensors(ctx, backend)
```

### Arguments

ctx	GGML context
backend	Backend handle

### Value

Backend buffer object

---

ggml\_backend\_buffer\_free  
*Free Backend Buffer*

---

### Description

Frees a backend buffer and all associated memory.

### Usage

```
ggml_backend_buffer_free(buffer)
```

### Arguments

buffer	Backend buffer object
--------	-----------------------

---

`ggml_backend_buffer_get_size`  
*Get Backend Buffer Size*

---

**Description**

Returns the total size of a backend buffer.

**Usage**

```
ggml_backend_buffer_get_size(buffer)
```

**Arguments**

buffer      Backend buffer object

**Value**

Size in bytes

---

`ggml_backend_buffer_name`  
*Get Backend Buffer Name*

---

**Description**

Returns the name/type of a backend buffer.

**Usage**

```
ggml_backend_buffer_name(buffer)
```

**Arguments**

buffer      Backend buffer object

**Value**

Character string with buffer name

---

ggml\_backend\_cpu\_init *Initialize CPU Backend*

---

### Description

Creates a new CPU backend instance for graph computation.

### Usage

```
ggml_backend_cpu_init()
```

### Value

Backend pointer

---

---

ggml\_backend\_cpu\_set\_n\_threads  
*Set CPU Backend Threads*

---

### Description

Sets the number of threads for CPU backend computation.

### Usage

```
ggml_backend_cpu_set_n_threads(backend, n_threads)
```

### Arguments

backend	CPU backend pointer
n_threads	Number of threads

### Value

NULL invisibly

---

`ggml_backend_free`      *Free Backend*

---

**Description**

Releases resources associated with a backend.

**Usage**

```
ggml_backend_free(backend)
```

**Arguments**

backend      Backend pointer

**Value**

NULL invisibly

---

`ggml_backend_graph_compute`  
                  *Compute Graph with Backend*

---

**Description**

Executes computation graph using specified backend.

**Usage**

```
ggml_backend_graph_compute(backend, graph)
```

**Arguments**

backend      Backend pointer  
graph        Graph pointer

**Value**

Status code (0 = success)

---

ggml_backend_name	<i>Get Backend Name</i>
-------------------	-------------------------

---

**Description**

Returns the name of the backend (e.g., "CPU").

**Usage**

```
ggml_backend_name(backend)
```

**Arguments**

backend	Backend pointer
---------	-----------------

**Value**

Character string name

---

ggml_backend_sched_alloc_graph	<i>Allocate graph on scheduler</i>
--------------------------------	------------------------------------

---

**Description**

Allocates memory for a graph across the scheduler's backends. Must be called before computing the graph.

**Usage**

```
ggml_backend_sched_alloc_graph(sched, graph)
```

**Arguments**

sched	Scheduler pointer
graph	Graph pointer

**Value**

Logical indicating success

---

`ggml_backend_sched_free`

*Free backend scheduler*

---

### Description

Releases resources associated with the backend scheduler.

### Usage

`ggml_backend_sched_free(sched)`

### Arguments

<code>sched</code>	Scheduler pointer from <code>ggml_backend_sched_new()</code>
--------------------	--

### Value

`NULL` (invisible)

### Examples

```
## Not run:
sched <- ggml_backend_sched_new(list(gpu_backend))
ggml_backend_sched_free(sched)

## End(Not run)
```

---

`ggml_backend_sched_get_backend`

*Get backend from scheduler*

---

### Description

Returns a specific backend from the scheduler by index.

### Usage

`ggml_backend_sched_get_backend(sched, index = 0L)`

### Arguments

<code>sched</code>	Scheduler pointer
<code>index</code>	Backend index (0-based)

### Value

Backend pointer

---

ggml\_backend\_sched\_get\_n\_backends  
*Get number of backends in scheduler*

---

**Description**

Returns the number of backends managed by the scheduler.

**Usage**

```
ggml_backend_sched_get_n_backends(sched)
```

**Arguments**

sched            Scheduler pointer

**Value**

Integer count of backends

---

ggml\_backend\_sched\_get\_n\_copies  
*Get number of tensor copies*

---

**Description**

Returns the number of tensor copies made in the last computed graph. Copies occur when data needs to be transferred between backends.

**Usage**

```
ggml_backend_sched_get_n_copies(sched)
```

**Arguments**

sched            Scheduler pointer

**Value**

Integer count of copies

---

```
ggml_backend_sched_get_n_splits
    Get number of graph splits
```

---

### Description

Returns the number of splits in the last computed graph. Higher numbers indicate more distribution across backends.

### Usage

```
ggml_backend_sched_get_n_splits(sched)
```

### Arguments

<code>sched</code>	Scheduler pointer
--------------------	-------------------

### Value

Integer count of splits

---

```
ggml_backend_sched_get_tensor_backend
    Get tensor backend assignment
```

---

### Description

Returns which backend a tensor is assigned to.

### Usage

```
ggml_backend_sched_get_tensor_backend(sched, tensor)
```

### Arguments

<code>sched</code>	Scheduler pointer
<code>tensor</code>	Tensor pointer

### Value

Backend pointer or NULL if not assigned

---

```
ggml_backend_sched_graph_compute
    Compute graph using scheduler
```

---

## Description

Computes a graph by distributing work across multiple backends. This is the main function for multi-GPU computation.

## Usage

```
ggml_backend_sched_graph_compute(sched, graph)
```

## Arguments

sched	Scheduler pointer
graph	Graph pointer

## Value

Status code (0 = success)

## Examples

```
## Not run:
# Multi-GPU example
if (ggml_vulkan_available() && ggml_vulkan_device_count() >= 2) {
  gpu1 <- ggml_vulkan_init(0)
  gpu2 <- ggml_vulkan_init(1)
  sched <- ggml_backend_sched_new(list(gpu1, gpu2))

  ctx <- ggml_init(64 * 1024 * 1024)
  a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10000)
  b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10000)
  ggml_set_f32(a, rnorm(10000))
  ggml_set_f32(b, rnorm(10000))

  c <- ggml_add(ctx, a, b)
  graph <- ggml_build_forward_expand(ctx, c)

  # Reserve memory
  ggml_backend_sched_reserve(sched, graph)

  # Compute using both GPUs
  ggml_backend_sched_graph_compute(sched, graph)

  result <- ggml_get_f32(c)

  cat("Splits:", ggml_backend_sched_get_n_splits(sched), "\n")
```

```

cat("Copies:", ggml_backend_sched_get_n_copies(sched), "\n")

ggml_free(ctx)
ggml_backend_sched_free(sched)
ggml_vulkan_free(gpu1)
ggml_vulkan_free(gpu2)
}

## End(Not run)

```

---

`ggml_backend_sched_graph_compute_async`  
*Compute graph asynchronously*

---

## Description

Computes a graph asynchronously across backends. Use `ggml_backend_sched_synchronize()` to wait for completion.

## Usage

```
ggml_backend_sched_graph_compute_async(sched, graph)
```

## Arguments

<code>sched</code>	Scheduler pointer
<code>graph</code>	Graph pointer

## Value

Status code (0 = success)

---

`ggml_backend_sched_new`  
*Create a new backend scheduler*

---

## Description

Creates a scheduler that can distribute computation across multiple backends (GPUs, CPU). A CPU backend is automatically added as a fallback. Backends with lower index have higher priority.

## Usage

```
ggml_backend_sched_new(backends, parallel = TRUE, graph_size = 2048)
```

**Arguments**

<code>backends</code>	List of backend pointers (from <code>ggml_vulkan_init()</code> or <code>ggml_backend_cpu_init()</code> ). Note: A CPU backend is automatically added, so you only need to specify GPU backends.
<code>parallel</code>	Logical, whether to run backends in parallel (default: TRUE)
<code>graph_size</code>	Expected maximum graph size (default: 2048)

**Value**

Scheduler pointer

**Examples**

```
## Not run:
if (ggml_vulkan_available() && ggml_vulkan_device_count() >= 2) {
  # Create two GPU backends (CPU is added automatically)
  gpu1 <- ggml_vulkan_init(0)
  gpu2 <- ggml_vulkan_init(1)

  # Create scheduler with both GPUs + CPU (automatic)
  sched <- ggml_backend_sched_new(list(gpu1, gpu2), parallel = TRUE)

  # The scheduler now has 3 backends: GPU1, GPU2, CPU
  cat("Backends:", ggml_backend_sched_get_n_backends(sched), "\n")

  # Use scheduler...

  # Cleanup
  ggml_backend_sched_free(sched)
  ggml_vulkan_free(gpu1)
  ggml_vulkan_free(gpu2)
}

## End(Not run)
```

**ggml\_backend\_sched\_reserve**

*Reserve memory for scheduler*

**Description**

Pre-allocates memory based on a measurement graph. This should be called before using the scheduler to compute graphs.

**Usage**

```
ggml_backend_sched_reserve(sched, graph)
```

**Arguments**

<code>sched</code>	Scheduler pointer
<code>graph</code>	Graph pointer to measure memory requirements

**Value**

Logical indicating success

**Examples**

```
## Not run:
sched <- ggml_backend_sched_new(list(gpu_backend))
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 1000)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 1000)
c <- ggml_add(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, c)

# Reserve memory based on this graph
ggml_backend_sched_reserve(sched, graph)

## End(Not run)
```

`ggml_backend_sched_reset`  
*Reset scheduler*

**Description**

Resets the scheduler, deallocating all tensors. Must be called before changing node backends or allocating a new graph.

**Usage**

```
ggml_backend_sched_reset(sched)
```

**Arguments**

<code>sched</code>	Scheduler pointer
--------------------	-------------------

**Value**

NULL (invisible)

---

```
ggml_backend_sched_set_tensor_backend  
    Set tensor backend assignment
```

---

**Description**

Manually assigns a specific tensor to run on a specific backend. This overrides automatic scheduling.

**Usage**

```
ggml_backend_sched_set_tensor_backend(sched, tensor, backend)
```

**Arguments**

sched	Scheduler pointer
tensor	Tensor pointer
backend	Backend pointer to assign tensor to

**Value**

NULL (invisible)

---

```
ggml_backend_sched_synchronize  
    Synchronize scheduler
```

---

**Description**

Waits for all asynchronous operations to complete.

**Usage**

```
ggml_backend_sched_synchronize(sched)
```

**Arguments**

sched	Scheduler pointer
-------	-------------------

**Value**

NULL (invisible)

---

```
ggml_backend_tensor_get_data
    Get Tensor Data via Backend
```

---

## Description

Gets tensor data using the backend API. This works with tensors allocated on any backend, not just CPU.

## Usage

```
ggml_backend_tensor_get_data(tensor, offset = 0, n_elements = NULL)
```

## Arguments

<code>tensor</code>	Tensor pointer
<code>offset</code>	Byte offset (default: 0)
<code>n_elements</code>	Number of elements to retrieve (NULL for all)

## Value

R vector with tensor data

---

```
ggml_backend_tensor_set_data
    Set Tensor Data via Backend
```

---

## Description

Sets tensor data using the backend API. This works with tensors allocated on any backend, not just CPU.

## Usage

```
ggml_backend_tensor_set_data(tensor, data, offset = 0)
```

## Arguments

<code>tensor</code>	Tensor pointer
<code>data</code>	R vector with data to set
<code>offset</code>	Byte offset (default: 0)

---

ggml\_build\_forward\_expand  
*Build forward expand*

---

## Description

Builds a computation graph from the output tensor, expanding backwards to include all dependencies.

Creates a computation graph by expanding backwards from the output tensor

## Usage

```
ggml_build_forward_expand(ctx, tensor)  
ggml_build_forward_expand(ctx, tensor)
```

## Arguments

ctx	GGML context
tensor	Output tensor of the computation

## Value

Graph pointer  
Graph object (external pointer)

## Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
ggml_set_f32(a, c(1, 2, 3, 4, 5))  
ggml_set_f32(b, c(5, 4, 3, 2, 1))  
result <- ggml_add(ctx, a, b)  
graph <- ggml_build_forward_expand(ctx, result)  
ggml_graph_compute(ctx, graph)  
ggml_free(ctx)  
  
## End(Not run)  
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
ggml_set_f32(a, 1:10)  
ggml_set_f32(b, 11:20)  
c <- ggml_add(ctx, a, b)  
graph <- ggml_build_forward_expand(ctx, c)
```

```
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(c)
ggml_free(ctx)

## End(Not run)
```

**ggml\_ceil***Ceiling (Graph)***Description**

Creates a graph node for element-wise ceiling: ceil(x)

**Usage**

```
ggml_ceil(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the ceil operation

**ggml\_clamp***Clamp (Graph)***Description**

Creates a graph node for clamping values to a range: clamp(x, min, max)

**Usage**

```
ggml_clamp(ctx, a, min_val, max_val)
```

**Arguments**

ctx	GGML context
a	Input tensor
min_val	Minimum value
max_val	Maximum value

**Value**

Tensor with values clamped to [min\_val, max\_val]

---

**ggml\_concat**      *Concatenate Tensors (Graph)*

---

**Description**

Concatenates two tensors along a specified dimension. CRITICAL for KV-cache operations in transformers.

**Usage**

```
ggml_concat(ctx, a, b, dim = 0)
```

**Arguments**

ctx	GGML context
a	First tensor
b	Second tensor (must match a in all dimensions except the concat dim)
dim	Dimension along which to concatenate (0-3)

**Value**

Concatenated tensor

**Examples**

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
a <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 4, 3)  
b <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 4, 2)  
ggml_set_f32(a, rnorm(12))  
ggml_set_f32(b, rnorm(8))  
# Concatenate along dimension 1: result is 4x5  
c <- ggml_concat(ctx, a, b, 1)  
graph <- ggml_build_forward_expand(ctx, c)  
ggml_graph_compute(ctx, graph)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml_cont	<i>Make Contiguous (Graph)</i>
-----------	--------------------------------

---

### Description

Makes a tensor contiguous in memory. Required after permute/transpose before some operations.

### Usage

```
ggml_cont(ctx, a)
```

### Arguments

ctx	GGML context
a	Input tensor

### Value

Contiguous tensor

---

ggml_conv_1d	<i>ID Convolution (Graph)</i>
--------------	-------------------------------

---

### Description

Applies 1D convolution to input data.

### Usage

```
ggml_conv_1d(ctx, a, b, s0 = 1L, p0 = 0L, d0 = 1L)
```

### Arguments

ctx	GGML context
a	Convolution kernel tensor
b	Input data tensor
s0	Stride (default 1)
p0	Padding (default 0)
d0	Dilation (default 1)

### Value

Convolved tensor

---

ggml_conv_2d	<i>2D Convolution (Graph)</i>
--------------	-------------------------------

---

## Description

Applies 2D convolution to input data.

## Usage

```
ggml_conv_2d(ctx, a, b, s0 = 1L, s1 = 1L, p0 = 0L, p1 = 0L, d0 = 1L, d1 = 1L)
```

## Arguments

ctx	GGML context
a	Convolution kernel tensor [KW, KH, IC, OC]
b	Input data tensor [W, H, C, N]
s0	Stride dimension 0 (default 1)
s1	Stride dimension 1 (default 1)
p0	Padding dimension 0 (default 0)
p1	Padding dimension 1 (default 0)
d0	Dilation dimension 0 (default 1)
d1	Dilation dimension 1 (default 1)

## Value

Convolved tensor

---

ggml_conv_transpose_1d	<i>Transposed 1D Convolution (Graph)</i>
------------------------	--

---

## Description

Applies transposed 1D convolution (deconvolution) to input data.

## Usage

```
ggml_conv_transpose_1d(ctx, a, b, s0 = 1L, p0 = 0L, d0 = 1L)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Convolution kernel tensor
<code>b</code>	Input data tensor
<code>s0</code>	Stride (default 1)
<code>p0</code>	Padding (default 0)
<code>d0</code>	Dilation (default 1)

**Value**

Transposed convolved tensor

`ggml_cos`*Cosine (Graph)***Description**

Creates a graph node for element-wise cosine:  $\cos(x)$

**Usage**

```
ggml_cos(ctx, a)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor

**Value**

Tensor representing the cos operation

ggml\_cpu\_add

*Element-wise Addition (CPU Direct)***Description**

Performs element-wise addition of two tensors using direct CPU computation. Returns the result as an R numeric vector. Does NOT use computation graphs.

**Usage**

```
ggml_cpu_add(a, b)
```

**Arguments**

- |   |  |
|---|--|
| a | First tensor (must be F32 type)                  |
| b | Second tensor (must be F32 type, same size as a) |

**Value**

Numeric vector containing the element-wise sum

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
ggml_set_f32(b, c(5, 4, 3, 2, 1))
ggml_cpu_add(a, b)
ggml_free(ctx)

## End(Not run)
```

ggml\_cpu\_mul

*Element-wise Multiplication (CPU Direct)***Description**

Performs element-wise multiplication of two tensors using direct CPU computation. Returns the result as an R numeric vector. Does NOT use computation graphs.

**Usage**

```
ggml_cpu_mul(a, b)
```

**Arguments**

- a                  First tensor (must be F32 type)
- b                  Second tensor (must be F32 type, same size as a)

**Value**

Numeric vector containing the element-wise product

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
ggml_set_f32(b, c(2, 2, 2, 2, 2))
ggml_cpu_mul(a, b)
ggml_free(ctx)

## End(Not run)
```

**ggml\_cpy***Copy Tensor with Type Conversion (Graph)***Description**

Copies tensor a into tensor b, performing type conversion if needed. The tensors must have the same number of elements. CRITICAL for type casting operations (e.g., F32 to F16).

**Usage**

```
ggml_cpy(ctx, a, b)
```

**Arguments**

- ctx                GGML context
- a                  Source tensor
- b                  Destination tensor (defines output type and shape)

**Value**

Tensor representing the copy operation (returns b with a's data)

## Examples

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
# Create F32 tensor  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 100)  
ggml_set_f32(a, rnorm(100))  
# Create F16 tensor for output  
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F16, 100)  
# Copy with F32 -> F16 conversion  
result <- ggml_cpy(ctx, a, b)  
graph <- ggml_build_forward_expand(ctx, result)  
ggml_graph_compute(ctx, graph)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_cycles

*Get CPU Cycles*

---

## Description

Returns the current CPU cycle count. Useful for low-level benchmarking.

## Usage

```
ggml_cycles()
```

## Value

Numeric value representing CPU cycles

## Examples

```
## Not run:  
ggml_cycles()  
  
## End(Not run)
```

`ggml_cycles_per_ms`      *Get CPU Cycles per Millisecond*

### Description

Returns an estimate of CPU cycles per millisecond. Useful for converting cycle counts to time.

### Usage

```
ggml_cycles_per_ms()
```

### Value

Numeric value representing cycles per millisecond

### Examples

```
## Not run:  
ggml_cycles_per_ms()  
  
## End(Not run)
```

`ggml_diag`      *Diagonal Matrix (Graph)*

### Description

Creates a diagonal matrix from a vector. For vector `a[n]`, produces matrix with `a` on the diagonal.

### Usage

```
ggml_diag(ctx, a)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input vector tensor

### Value

Diagonal matrix tensor

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 3)
ggml_set_f32(a, c(1, 2, 3))
d <- ggml_diag(ctx, a) # 3x3 diagonal matrix
graph <- ggml_build_forward_expand(ctx, d)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

**ggml\_diag\_mask\_inf**      *Diagonal Mask with -Inf (Graph)*

## Description

Creates a graph node that sets elements above the diagonal to -Inf. This is used for causal (autoregressive) attention masking.

## Usage

```
ggml_diag_mask_inf(ctx, a, n_past)
```

## Arguments

ctx	GGML context
a	Input tensor (typically attention scores)
n_past	Number of past tokens (shifts the diagonal). Use 0 for standard causal masking where position i can only attend to positions $\leq i$ .

## Details

In causal attention, we want each position to only attend to itself and previous positions. Setting future positions to -Inf ensures that after softmax, they contribute 0 attention weight.

The n\_past parameter allows for KV-cache scenarios where the diagonal needs to be shifted to account for previously processed tokens.

## Value

Tensor with same shape as input, elements above diagonal set to -Inf

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
# Create attention scores matrix
scores <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 4, 4)
ggml_set_f32(scores, rep(1, 16))
# Apply causal mask
masked <- ggml_diag_mask_inf(ctx, scores, 0)
graph <- ggml_build_forward_expand(ctx, masked)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

### *ggml\_diag\_mask\_inf\_inplace*

*Diagonal Mask with -Inf In-place (Graph)*

## Description

In-place version of `ggml_diag_mask_inf`. Returns a view of the input tensor.

## Usage

```
ggml_diag_mask_inf_inplace(ctx, a, n_past)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (will be modified in-place)
<code>n_past</code>	Number of past tokens

## Value

View of input tensor with elements above diagonal set to -Inf

---

ggml\_diag\_mask\_zero     *Diagonal Mask with Zero (Graph)*

---

### Description

Creates a graph node that sets elements above the diagonal to 0. Alternative to -Inf masking for certain use cases.

### Usage

```
ggml_diag_mask_zero(ctx, a, n_past)
```

### Arguments

ctx	GGML context
a	Input tensor
n_past	Number of past tokens

### Value

Tensor with same shape as input, elements above diagonal set to 0

---

ggml\_div                *Element-wise Division (Graph)*

---

### Description

Creates a graph node for element-wise division.

### Usage

```
ggml_div(ctx, a, b)
```

### Arguments

ctx	GGML context
a	First tensor (numerator)
b	Second tensor (denominator, same shape as a)

### Value

Tensor representing the division operation (a / b)

## Examples

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(10, 20, 30, 40, 50))
ggml_set_f32(b, c(2, 2, 2, 2, 2))
result <- ggml_div(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

**ggml\_dup**

*Duplicate Tensor (Graph)*

## Description

Creates a graph node that copies a tensor. This is a graph operation that must be computed using `ggml_build_forward_expand()` and `ggml_graph_compute()`. Unlike `ggml_dup_tensor` which just allocates, this creates a copy operation in the graph.

## Usage

```
ggml_dup(ctx, a)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor

## Value

Tensor representing the copy operation

## Examples

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
b <- ggml_dup(ctx, a)
graph <- ggml_build_forward_expand(ctx, b)
ggml_graph_compute(ctx, graph)
ggml_get_f32(b)
ggml_free(ctx)

## End(Not run)
```

---

ggml_dup_tensor	<i>Duplicate Tensor</i>
-----------------	-------------------------

---

### Description

Creates a copy of a tensor with the same shape and type

### Usage

```
ggml_dup_tensor(ctx, tensor)
```

### Arguments

ctx	GGML context
tensor	Tensor to duplicate

### Value

New tensor pointer with same shape

### Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 100)  
b <- ggml_dup_tensor(ctx, a)  
ggml_nelements(b)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml_element_size	<i>Get Element Size</i>
-------------------	-------------------------

---

### Description

Returns the size of a single element in the tensor.

### Usage

```
ggml_element_size(tensor)
```

### Arguments

tensor	Tensor pointer
--------	----------------

**Value**

Element size in bytes

---

`ggml_elu`

*ELU Activation (Graph)*

---

**Description**

Creates a graph node for ELU (Exponential Linear Unit) activation.  $\text{ELU}(x) = x \text{ if } x > 0, \text{ else } \alpha * (\exp(x) - 1)$  where  $\alpha = 1$ .

**Usage**

```
ggml_elu(ctx, a)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor

**Value**

Tensor representing the ELU operation

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-2, -1, 0, 1, 2))
r <- ggml_elu(ctx, a)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r)
ggml_free(ctx)

## End(Not run)
```

---

ggml\_estimate\_memory    *Estimate Required Memory*

---

### Description

Helper function to estimate memory needed for a tensor

### Usage

```
ggml_estimate_memory(type = GGML_TYPE_F32, ne0, ne1 = 1, ne2 = 1, ne3 = 1)
```

### Arguments

type	Tensor type (GGML_TYPE_F32, etc)
ne0	Size of dimension 0
ne1	Size of dimension 1 (optional)
ne2	Size of dimension 2 (optional)
ne3	Size of dimension 3 (optional)

### Value

Estimated memory in bytes

### Examples

```
## Not run:  
# For 1000x1000 F32 matrix  
ggml_estimate_memory(GGML_TYPE_F32, 1000, 1000)  
  
## End(Not run)
```

---

ggml\_exp

*Exponential (Graph)*

---

### Description

Creates a graph node for element-wise exponential:  $\exp(x)$

### Usage

```
ggml_exp(ctx, a)
```

### Arguments

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the exp operation

---

`ggml_flash_attn_back`    *Flash Attention Backward (Graph)*

---

**Description**

Backward pass for Flash Attention. Used during training to compute gradients through attention.

**Usage**

```
ggml_flash_attn_back(ctx, q, k, v, d, masked = TRUE)
```

**Arguments**

<code>ctx</code>	GGML context
<code>q</code>	Query tensor (same as forward pass)
<code>k</code>	Key tensor (same as forward pass)
<code>v</code>	Value tensor (same as forward pass)
<code>d</code>	Gradient tensor from upstream (same shape as forward output)
<code>masked</code>	Logical: whether causal masking was used in forward pass

**Value**

Gradient tensor

---

`ggml_flash_attn_ext`    *Flash Attention (Graph)*

---

**Description**

Creates a graph node for Flash Attention computation. This is a memory-efficient implementation of scaled dot-product attention.

**Usage**

```
ggml_flash_attn_ext(
    ctx,
    q,
    k,
    v,
    mask = NULL,
    scale,
    max_bias = 0,
    logit_softcap = 0
)
```

## Arguments

ctx	GGML context
q	Query tensor of shape [head_dim, n_head, n_tokens, batch]
k	Key tensor of shape [head_dim, n_head_kv, n_kv, batch]
v	Value tensor of shape [head_dim, n_head_kv, n_kv, batch]
mask	Optional attention mask tensor (NULL for no mask). For causal attention, use ggml_diag_mask_inf instead.
scale	Attention scale factor, typically $1/\sqrt{\text{head\_dim}}$
max_bias	Maximum ALiBi bias (0.0 to disable ALiBi)
logit_softcap	Logit soft-capping value (0.0 to disable). Used by some models like Gemma 2.

## Details

Flash Attention computes:  $\text{softmax}(Q * K^T / \text{scale} + \text{mask}) * V$

Key features:

- Memory efficient:  $O(n)$  instead of  $O(n^2)$  memory for attention matrix
- Supports grouped-query attention (GQA) when  $n_{\text{head\_kv}} < n_{\text{head}}$
- Supports multi-query attention (MQA) when  $n_{\text{head\_kv}} = 1$
- Optional ALiBi (Attention with Linear Biases) for position encoding
- Optional logit soft-capping for numerical stability

## Value

Attention output tensor of shape [head\_dim, n\_head, n\_tokens, batch]

## Examples

```
## Not run:
ctx <- ggml_init(64 * 1024 * 1024)
head_dim <- 64
n_head <- 8
n_head_kv <- 2 # GQA with 4:1 ratio
seq_len <- 32
q <- ggml_new_tensor_4d(ctx, GGML_TYPE_F32, head_dim, n_head, seq_len, 1)
k <- ggml_new_tensor_4d(ctx, GGML_TYPE_F32, head_dim, n_head_kv, seq_len, 1)
v <- ggml_new_tensor_4d(ctx, GGML_TYPE_F32, head_dim, n_head_kv, seq_len, 1)
ggml_set_f32(q, rnorm(head_dim * n_head * seq_len))
ggml_set_f32(k, rnorm(head_dim * n_head_kv * seq_len))
ggml_set_f32(v, rnorm(head_dim * n_head_kv * seq_len))
# Scale = 1/sqrt(head_dim)
scale <- 1.0 / sqrt(head_dim)
# Compute attention
out <- ggml_flash_attn_ext(ctx, q, k, v, NULL, scale, 0.0, 0.0)
graph <- ggml_build_forward_expand(ctx, out)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

`ggml_floor`                    *Floor (Graph)*

### Description

Creates a graph node for element-wise floor:  $\text{floor}(x)$

### Usage

```
ggml_floor(ctx, a)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor

### Value

Tensor representing the floor operation

`ggml_free`                    *Free GGML context*

### Description

Free GGML context

### Usage

```
ggml_free(ctx)
```

### Arguments

<code>ctx</code>	Context pointer
------------------	-----------------

### Value

NULL (invisible)

### Examples

```
## Not run:
ctx <- ggml_init(1024 * 1024)
ggml_free(ctx)

## End(Not run)
```

---

ggml\_gallocr\_alloc\_graph  
*Allocate Memory for Graph*

---

**Description**

Allocates memory for all tensors in the computation graph. This must be called before computing the graph.

**Usage**

```
ggml_gallocr_alloc_graph(galloc, graph)
```

**Arguments**

galloc	Graph allocator object
graph	Graph object

**Value**

TRUE on success, FALSE on failure

**Examples**

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
galloc <- ggml_gallocr_new()  
  
# Create graph  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
b <- ggml_relu(ctx, a)  
graph <- ggml_build_forward_expand(ctx, b)  
  
# Allocate and compute  
ggml_gallocr_alloc_graph(galloc, graph)  
ggml_graph_compute(ctx, graph)  
  
ggml_gallocr_free(galloc)  
ggml_free(ctx)  
  
## End(Not run)
```

---

`ggml_gallocr_free`      *Free Graph Allocator*

---

## Description

Frees a graph allocator and all associated buffers.

## Usage

```
ggml_gallocr_free(galloc)
```

## Arguments

`galloc`      Graph allocator object

---

`ggml_gallocr_get_buffer_size`  
*Get Graph Allocator Buffer Size*

---

## Description

Returns the size of the buffer used by the graph allocator.

## Usage

```
ggml_gallocr_get_buffer_size(galloc, buffer_id = 0L)
```

## Arguments

`galloc`      Graph allocator object  
`buffer_id`      Buffer ID (default: 0 for single-buffer allocator)

## Value

Size in bytes

---

ggml\_gallocr\_new      *Create Graph Allocator*

---

## Description

Creates a new graph allocator for efficient memory management. The allocator can automatically allocate and reuse memory for graph tensors.

## Usage

```
ggml_gallocr_new()
```

## Value

Graph allocator object (external pointer)

## Examples

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
galloc <- ggml_gallocr_new()  
  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
b <- ggml_relu(ctx, a)  
graph <- ggml_build_forward_expand(ctx, b)  
  
# Allocate graph  
ggml_gallocr_alloc_graph(galloc, graph)  
  
ggml_gallocr_free(galloc)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_gallocr\_reserve    *Reserve Memory for Graph*

---

## Description

Pre-allocates memory for a graph. This is optional but recommended when running the same graph multiple times to avoid reallocation.

## Usage

```
ggml_gallocr_reserve(galloc, graph)
```

**Arguments**

<code>galloc</code>	Graph allocator object
<code>graph</code>	Graph object

**Value**

TRUE on success, FALSE on failure

`ggml_geglu`

*GeGLU (GELU Gated Linear Unit) (Graph)*

**Description**

Creates a graph node for GeGLU operation. GeGLU uses GELU as the activation function on the first half. CRITICAL for models like GPT-NeoX and Falcon.

**Usage**

```
ggml_geglu(ctx, a)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (first dimension must be even)

**Details**

Formula: output = GELU(x) \* gate

**Value**

Tensor with half the first dimension of input

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 8, 3)
ggml_set_f32(a, rnorm(24))
r <- ggml_geglu(ctx, a)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r) # Shape: 4x3
ggml_free(ctx)

## End(Not run)
```

---

ggml_geglu_quick	<i>GeGLU Quick (Fast GeGLU) (Graph)</i>
------------------	---

---

### Description

Creates a graph node for fast GeGLU approximation. Uses faster but less accurate GELU approximation for gating.

### Usage

```
ggml_geglu_quick(ctx, a)
```

### Arguments

ctx	GGML context
a	Input tensor (first dimension must be even)

### Value

Tensor with half the first dimension of input

---

ggml_geglu_split	<i>GeGLU Split (Graph)</i>
------------------	----------------------------

---

### Description

Creates a graph node for GeGLU with separate input and gate tensors.

### Usage

```
ggml_geglu_split(ctx, a, b)
```

### Arguments

ctx	GGML context
a	Input tensor (the values to be gated)
b	Gate tensor (same shape as a)

### Details

Formula: output = GELU(a) \* b

### Value

Tensor with same shape as input tensors

---

**ggml\_gelu***GELU Activation (Graph)*

---

**Description**

Creates a graph node for GELU (Gaussian Error Linear Unit) activation. CRITICAL for GPT models.

**Usage**

```
ggml_gelu(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the GELU operation

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-2, -1, 0, 1, 2))
result <- ggml_gelu(ctx, a)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

---

**ggml\_gelu\_erf***Exact GELU Activation (Graph)*

---

**Description**

Creates a graph node for exact GELU using the error function (erf).  $\text{GELU}(x) = x * 0.5 * (1 + \text{erf}(x / \sqrt{2}))$ . More accurate than approximate GELU but potentially slower on some backends.

**Usage**

```
ggml_gelu_erf(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the exact GELU operation

**Examples**

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
ggml_set_f32(a, c(-2, -1, 0, 1, 2))  
r <- ggml_gelu_erf(ctx, a)  
graph <- ggml_build_forward_expand(ctx, r)  
ggml_graph_compute(ctx, graph)  
result <- ggml_get_f32(r)  
ggml_free(ctx)  
  
## End(Not run)
```

---

**ggml\_gelu\_quick**      *GELU Quick Activation (Graph)*

---

**Description**

Creates a graph node for fast approximation of GELU. Faster than standard GELU with minimal accuracy loss.

**Usage**

```
ggml_gelu_quick(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the GELU quick operation

---

`ggml_get_f32`      *Get F32 data*

---

## Description

Get F32 data

Get F32 Data

## Usage

```
ggml_get_f32(tensor)  
ggml_get_f32(tensor)
```

## Arguments

`tensor`      Tensor

## Value

Numeric vector with tensor values

Numeric vector

## Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
tensor <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
ggml_set_f32(tensor, c(1, 2, 3, 4, 5))  
ggml_get_f32(tensor)  
ggml_free(ctx)  
  
## End(Not run)  
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
ggml_set_f32(t, c(1, 2, 3, 4, 5))  
ggml_get_f32(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

`ggml_get_i32`*Get I32 Data*

---

**Description**

Gets integer data from an I32 tensor (e.g., from ggml\_argmax)

**Usage**

```
ggml_get_i32(tensor)
```

**Arguments**

`tensor` Tensor of type GGML\_TYPE\_I32

**Value**

Integer vector

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
pos <- ggml_new_tensor_1d(ctx, GGML_TYPE_I32, 10)  
ggml_set_i32(pos, 0:9)  
ggml_get_i32(pos)  
ggml_free(ctx)  
  
## End(Not run)
```

---

`ggml_get_max_tensor_size`*Get Maximum Tensor Size*

---

**Description**

Returns the maximum tensor size that can be allocated in the context

**Usage**

```
ggml_get_max_tensor_size(ctx)
```

**Arguments**

`ctx` GGML context

**Value**

Maximum tensor size in bytes

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
ggml_get_max_tensor_size(ctx)  
ggml_free(ctx)  
  
## End(Not run)
```

---

`ggml_get_mem_size`      *Get Context Memory Size*

---

**Description**

Returns the total memory pool size of the context

**Usage**

```
ggml_get_mem_size(ctx)
```

**Arguments**

`ctx`      GGML context

**Value**

Total memory size in bytes

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
ggml_get_mem_size(ctx)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml_get_name	<i>Get Tensor Name</i>
---------------	------------------------

---

### Description

Retrieves the name of a tensor.

### Usage

```
ggml_get_name(tensor)
```

### Arguments

tensor	Tensor pointer
--------	----------------

### Value

Character string name or NULL if not set

---

ggml_get_no_alloc	<i>Get No Allocation Mode</i>
-------------------	-------------------------------

---

### Description

Check if no-allocation mode is enabled

### Usage

```
ggml_get_no_alloc(ctx)
```

### Arguments

ctx	GGML context
-----	--------------

### Value

Logical indicating if no\_alloc is enabled

### Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
ggml_get_no_alloc(ctx)  
ggml_free(ctx)  
  
## End(Not run)
```

`ggml_get_n_threads`      *Get Number of Threads*

### Description

Get the current number of threads for GGML operations

### Usage

```
ggml_get_n_threads()
```

### Value

Number of threads

### Examples

```
## Not run:  
ggml_get_n_threads()  
  
## End(Not run)
```

`ggml_get_rows`      *Get Rows by Indices (Graph)*

### Description

Creates a graph node that extracts rows from a tensor by index. This is commonly used for embedding lookup in LLMs.

### Usage

```
ggml_get_rows(ctx, a, b)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Data tensor of shape [n_embd, n_rows, ...] - the embedding table
<code>b</code>	Index tensor (int32) of shape [n_indices] - which rows to extract

### Details

This operation is fundamental for embedding lookup in transformers: given a vocabulary embedding matrix and token indices, it retrieves the corresponding embedding vectors.

**Value**

Tensor of shape [n\_embd, n\_indices, ...] containing the selected rows

**Examples**

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
# Create embedding matrix: 10 tokens, 4-dim embeddings  
embeddings <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 4, 10)  
ggml_set_f32(embeddings, rnorm(40))  
# Token indices to look up  
indices <- ggml_new_tensor_1d(ctx, GGML_TYPE_I32, 3)  
ggml_set_i32(indices, c(0L, 5L, 2L))  
# Get embeddings for tokens 0, 5, 2  
result <- ggml_get_rows(ctx, embeddings, indices)  
graph <- ggml_build_forward_expand(ctx, result)  
ggml_graph_compute(ctx, graph)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_get\_rows\_back      *Get Rows Backward (Graph)*

---

**Description**

Backward pass for ggml\_get\_rows operation. Accumulates gradients at the original row positions.

**Usage**

```
ggml_get_rows_back(ctx, a, b, c)
```

**Arguments**

ctx	GGML context
a	Gradient of get_rows output
b	Index tensor (same as forward pass)
c	Reference tensor defining output shape

**Value**

Gradient tensor for the embedding matrix

**ggml\_glu***Generic GLU (Gated Linear Unit) (Graph)*

## Description

Creates a graph node for GLU operation with specified gating type. GLU splits the input tensor in half along the first dimension, applies an activation to the first half (x), and multiplies it with the second half (gate).

## Usage

```
ggml_glu(ctx, a, op, swapped = FALSE)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (first dimension must be even)
<code>op</code>	GLU operation type (GGML_GLU_OP_REGLU, GGML_GLU_OP_GEGLU, etc.)
<code>swapped</code>	If TRUE, swap x and gate halves (default FALSE)

## Details

Formula: output = activation(x) \* gate where x and gate are the two halves of the input tensor.

## Value

Tensor with shape [n/2, ...] where n is the first dimension of input

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
# Create tensor with 10 columns (will be split into 5 + 5)
a <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 4)
ggml_set_f32(a, rnorm(40))
# Apply SwiGLU
r <- ggml_glu(ctx, a, GGML_GLU_OP_SWIGLU, FALSE)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r) # Shape: 5x4
ggml_free(ctx)

## End(Not run)
```

---

GGML\_GLU\_OP\_REGLU      *GLU Operation Types*

---

### Description

Constants for GLU (Gated Linear Unit) operation types. Used with `ggml_glu()` and `ggml_glu_split()`.

### Usage

```
GGML_GLU_OP_REGLU  
GGML_GLU_OP_GEGLU  
GGML_GLU_OP_SWIGLU  
GGML_GLU_OP_SWIGLU_OAI  
GGML_GLU_OP_GEGLU_ERF  
GGML_GLU_OP_GEGLU_QUICK
```

### Format

Integer constants

An object of class `integer` of length 1.

### Details

- `GGML_GLU_OP_REGLU`: ReGLU - ReLU gating
- `GGML_GLU_OP_GEGLU`: GeGLU - GELU gating (used in GPT-NeoX, etc.)
- `GGML_GLU_OP_SWIGLU`: SwiGLU - SiLU/Swish gating (used in LLaMA)
- `GGML_GLU_OP_GEGLU_QUICK`: GeGLU with fast approximation

`ggml_glu_split`      *Generic GLU Split (Graph)*

## Description

Creates a graph node for GLU with separate input and gate tensors. Unlike standard GLU which splits a single tensor, this takes two separate tensors.

## Usage

```
ggml_glu_split(ctx, a, b, op)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (the values to be gated)
<code>b</code>	Gate tensor (same shape as a)
<code>op</code>	GLU operation type (GGML_GLU_OP_REGLU, GGML_GLU_OP_GEGLU, etc.)

## Value

Tensor with same shape as input tensors

`ggml_graph_compute`      *Compute graph*

## Description

Executes all operations in the computation graph.

Executes the computation graph using CPU backend

## Usage

```
ggml_graph_compute(ctx, graph)
```

```
ggml_graph_compute(ctx, graph)
```

## Arguments

<code>ctx</code>	GGML context
<code>graph</code>	Graph object created by <code>ggml_build_forward_expand</code>

**Value**

NULL (invisible)

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
result <- ggml_relu(ctx, a)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)
ggml_set_f32(a, 1:10)
ggml_set_f32(b, 11:20)
c <- ggml_add(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, c)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(c)
ggml_free(ctx)

## End(Not run)
```

**ggml\_graph\_compute\_with\_ctx**

*Compute Graph with Context (Alternative Method)*

**Description**

Computes the computation graph using the context-based method. This is an alternative to `ggml_graph_compute()` that uses `ggml_graph_plan()` and `ggml_graph_compute()` internally.

**Usage**

```
ggml_graph_compute_with_ctx(ctx, graph, n_threads = 0L)
```

**Arguments**

ctx	GGML context
graph	Graph object created by <code>ggml_build_forward_expand</code>
n_threads	Number of threads to use (0 for auto-detect, default: 0)

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)
ggml_set_f32(a, 1:10)
c <- ggml_relu(ctx, a)
graph <- ggml_build_forward_expand(ctx, c)
ggml_graph_compute_with_ctx(ctx, graph)
result <- ggml_get_f32(c)
ggml_free(ctx)

## End(Not run)
```

`ggml_graph_dump_dot`    *Export Graph to DOT Format*

## Description

Exports the computation graph to a DOT file for visualization. The DOT file can be converted to an image using Graphviz tools.

## Usage

```
ggml_graph_dump_dot(graph, leafs = NULL, filename)
```

## Arguments

<code>graph</code>	Graph object
<code>leafs</code>	Optional graph with leaf tensors (NULL for none)
<code>filename</code>	Output filename (should end with .dot)

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)
b <- ggml_relu(ctx, a)
graph <- ggml_build_forward_expand(ctx, b)
ggml_graph_dump_dot(graph, NULL, tempfile(fileext = ".dot"))
ggml_free(ctx)

## End(Not run)
```

---

**ggml\_graph\_get\_tensor** *Get Tensor from Graph by Name*

---

**Description**

Finds a tensor in the computation graph by its name

**Usage**

```
ggml_graph_get_tensor(graph, name)
```

**Arguments**

graph	Graph object
name	Character string with tensor name

**Value**

Tensor pointer or NULL if not found

---

**ggml\_graph\_node** *Get Graph Node*

---

**Description**

Gets a specific node (tensor) from the computation graph by index

**Usage**

```
ggml_graph_node(graph, i)
```

**Arguments**

graph	Graph object
i	Node index (0-based, negative indices count from end)

**Value**

Tensor pointer

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)
b <- ggml_add(ctx, a, a)
graph <- ggml_build_forward_expand(ctx, b)
# Get the last node (output)
output <- ggml_graph_node(graph, -1)
ggml_free(ctx)

## End(Not run)
```

**ggml\_graph\_n\_nodes**      *Get Number of Nodes in Graph*

## Description

Returns the number of computation nodes in the graph

## Usage

```
ggml_graph_n_nodes(graph)
```

## Arguments

graph	Graph object
-------	--------------

## Value

Integer number of nodes

**ggml\_graph\_overhead**      *Get Graph Overhead*

## Description

Returns the memory overhead required for a computation graph

## Usage

```
ggml_graph_overhead()
```

## Value

Size in bytes

---

ggml\_graph\_print      *Print Graph Information*

---

## Description

Prints debug information about the computation graph

## Usage

```
ggml_graph_print(graph)
```

## Arguments

graph	Graph object
-------	--------------

---

ggml\_graph\_reset      *Reset Graph (for backpropagation)*

---

## Description

Resets the computation graph for a new backward pass. NOTE: This function requires the graph to have gradients allocated (used for training/backpropagation). For inference-only graphs, this function will cause an error.

## Usage

```
ggml_graph_reset(graph)
```

## Arguments

graph	Graph object with gradients allocated
-------	---------------------------------------

`ggml_group_norm`      *Group Normalization (Graph)*

### Description

Creates a graph node for group normalization. Normalizes along  $ne0*ne1*n\_groups$  dimensions.  
Used in Stable Diffusion and other image generation models.

### Usage

```
ggml_group_norm(ctx, a, n_groups, eps = 1e-05)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor
<code>n_groups</code>	Number of groups to divide channels into
<code>eps</code>	Epsilon for numerical stability (default 1e-5)

### Value

Tensor representing the group norm operation

`ggml_group_norm_inplace`      *Group Normalization In-place (Graph)*

### Description

Creates a graph node for in-place group normalization.

### Usage

```
ggml_group_norm_inplace(ctx, a, n_groups, eps = 1e-05)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (will be modified in-place)
<code>n_groups</code>	Number of groups
<code>eps</code>	Epsilon for numerical stability (default 1e-5)

### Value

View of input tensor with group norm applied

<code>ggml_hardsigmoid</code>	<i>Hard Sigmoid Activation (Graph)</i>
-------------------------------	--

### Description

Creates a graph node for Hard Sigmoid activation.  $\text{HardSigmoid}(x) = \text{ReLU6}(x + 3) / 6 = \min(\max(0, x + 3), 6) / 6$ . A computationally efficient approximation of the sigmoid function.

### Usage

```
ggml_hardsigmoid(ctx, a)
```

### Arguments

ctx	GGML context
a	Input tensor

### Value

Tensor representing the Hard Sigmoid operation

### Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-4, -1, 0, 1, 4))
r <- ggml_hardsigmoid(ctx, a)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r) # [0, 0.333, 0.5, 0.667, 1]
ggml_free(ctx)

## End(Not run)
```

<code>ggml_hardswish</code>	<i>Hard Swish Activation (Graph)</i>
-----------------------------	--------------------------------------

### Description

Creates a graph node for Hard Swish activation.  $\text{HardSwish}(x) = x * \text{ReLU6}(x + 3) / 6 = x * \min(\max(0, x + 3), 6) / 6$ . Used in MobileNetV3 and other efficient architectures.

### Usage

```
ggml_hardswish(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the Hard Swish operation

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-4, -1, 0, 1, 4))
r <- ggml_hardswish(ctx, a)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r)
ggml_free(ctx)

## End(Not run)
```

**ggml\_im2col**

*Image to Column (Graph)*

**Description**

Transforms image data into column format for efficient convolution. This is a low-level operation used internally by convolution implementations.

**Usage**

```
ggml_im2col(
  ctx,
  a,
  b,
  s0,
  s1,
  p0,
  p1,
  d0,
  d1,
  is_2D = TRUE,
  dst_type = GGML_TYPE_F16
)
```

**Arguments**

ctx	GGML context
a	Convolution kernel tensor
b	Input data tensor
s0	Stride dimension 0
s1	Stride dimension 1
p0	Padding dimension 0
p1	Padding dimension 1
d0	Dilation dimension 0
d1	Dilation dimension 1
is_2D	Whether this is a 2D operation (default TRUE)
dst_type	Output type (default GGML_TYPE_F16)

**Value**

Transformed tensor in column format

---

ggml_init	<i>Initialize GGML context</i>
-----------	--------------------------------

---

**Description**

Initialize GGML context

**Usage**

```
ggml_init(mem_size = 16 * 1024 * 1024, no_alloc = FALSE)
```

**Arguments**

mem_size	Memory size in bytes
no_alloc	If TRUE, don't allocate memory for tensors (default: FALSE)

**Value**

GGML context pointer

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
ggml_free(ctx)  
  
## End(Not run)
```

`ggml_init_auto`      *Create Context with Auto-sizing*

### Description

Creates a context with automatically calculated size based on planned tensors

### Usage

```
ggml_init_auto(..., extra_mb = 10, type = GGML_TYPE_F32, no_alloc = FALSE)
```

### Arguments

...	Named arguments with tensor dimensions
extra_mb	Extra megabytes to add (default: 10)
type	Tensor type (default: GGML_TYPE_F32)
no_alloc	If TRUE, don't allocate memory for tensors (default: FALSE)

### Value

GGML context

### Examples

```
## Not run:
# For two 1000x1000 matrices
ctx <- ggml_init_auto(mat1 = c(1000, 1000), mat2 = c(1000, 1000))
ggml_free(ctx)

## End(Not run)
```

`ggml_is_available`      *Check if GGML is available*

### Description

Check if GGML is available

### Usage

```
ggml_is_available()
```

### Value

TRUE if GGML library is loaded

**Examples**

```
## Not run:  
ggml_is_available()  
  
## End(Not run)
```

---

ggml\_is\_contiguous      *Check if Tensor is Contiguous*

---

**Description**

Returns TRUE if tensor data is stored contiguously in memory

**Usage**

```
ggml_is_contiguous(tensor)
```

**Arguments**

tensor      Tensor pointer

**Value**

Logical

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
ggml_is_contiguous(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_is\_permuted      *Check if Tensor is Permuted*

---

**Description**

Returns TRUE if tensor dimensions have been permuted

**Usage**

```
ggml_is_permuted(tensor)
```

**Arguments**

<code>tensor</code>	Tensor pointer
---------------------	----------------

**Value**

Logical

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
t <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 20)
ggml_is_permuted(t)
ggml_free(ctx)

## End(Not run)
```

<code>ggml_is_transposed</code>	<i>Check if Tensor is Transposed</i>
---------------------------------	--------------------------------------

**Description**

Returns TRUE if tensor has been transposed

**Usage**

```
ggml_is_transposed(tensor)
```

**Arguments**

<code>tensor</code>	Tensor pointer
---------------------	----------------

**Value**

Logical

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
t <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 20)
ggml_is_transposed(t)
ggml_free(ctx)

## End(Not run)
```

---

ggml_l2_norm	<i>L2 Normalization (Graph)</i>
--------------	---------------------------------

---

### Description

Creates a graph node for L2 normalization (unit norm). Normalizes vectors to unit length:  $x / \|x\|_2$ . Used in RWKV v7 and embedding normalization.

### Usage

```
ggml_l2_norm(ctx, a, eps = 1e-05)
```

### Arguments

ctx	GGML context
a	Input tensor
eps	Epsilon for numerical stability (default 1e-5)

### Value

Tensor representing the L2 norm operation

---

ggml_l2_norm_inplace	<i>L2 Normalization In-place (Graph)</i>
----------------------	--

---

### Description

Creates a graph node for in-place L2 normalization.

### Usage

```
ggml_l2_norm_inplace(ctx, a, eps = 1e-05)
```

### Arguments

ctx	GGML context
a	Input tensor (will be modified in-place)
eps	Epsilon for numerical stability (default 1e-5)

### Value

View of input tensor with L2 norm applied

`ggml_leaky_relu`      *Leaky ReLU Activation (Graph)*

### Description

Creates a graph node for Leaky ReLU activation.  $\text{LeakyReLU}(x) = x \text{ if } x > 0, \text{ else } \text{negative\_slope} * x$ . Unlike standard ReLU, Leaky ReLU allows a small gradient for negative values.

### Usage

```
ggml_leaky_relu(ctx, a, negative_slope = 0.01, inplace = FALSE)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor
<code>negative_slope</code>	Slope for negative values (default: 0.01)
<code>inplace</code>	If TRUE, operation is performed in-place (default: FALSE)

### Value

Tensor representing the Leaky ReLU operation

### Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-2, -1, 0, 1, 2))
r <- ggml_leaky_relu(ctx, a, negative_slope = 0.1)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r) # [-0.2, -0.1, 0, 1, 2]
ggml_free(ctx)

## End(Not run)
```

`ggml_log`      *Natural Logarithm (Graph)*

### Description

Creates a graph node for element-wise natural logarithm:  $\log(x)$

**Usage**

```
ggml_log(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the log operation

---

**ggml\_mean***Mean (Graph)*

---

**Description**

Creates a graph node that computes the mean of all elements.

**Usage**

```
ggml_mean(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Scalar tensor with the mean

---

**ggml\_mul***Multiply tensors*

---

**Description**

Creates a graph node for element-wise multiplication.

**Usage**

```
ggml_mul(ctx, a, b)
```

```
ggml_mul(ctx, a, b)
```

**Arguments**

ctx	GGML context
a	First tensor
b	Second tensor (same shape as a)

**Value**

Tensor representing the multiplication operation  
 Tensor representing the multiplication operation

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
ggml_set_f32(b, c(2, 2, 2, 2, 2))
result <- ggml_mul(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(1, 2, 3, 4, 5))
ggml_set_f32(b, c(2, 2, 2, 2, 2))
result <- ggml_mul(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

**Description**

Creates a graph node for matrix multiplication. CRITICAL for LLM operations. For matrices A (m x n) and B (n x p), computes C = A \* B (m x p).

**Usage**

```
ggml_mul_mat(ctx, a, b)
```

**Arguments**

ctx	GGML context
a	First matrix tensor
b	Second matrix tensor

**Value**

Tensor representing the matrix multiplication

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
A <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 4, 3)
B <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 4, 2)
ggml_set_f32(A, 1:12)
ggml_set_f32(B, 1:8)
C <- ggml_mul_mat(ctx, A, B)
graph <- ggml_build_forward_expand(ctx, C)
ggml_graph_compute(ctx, graph)
ggml_get_f32(C)
ggml_free(ctx)

## End(Not run)
```

ggml\_mul\_mat\_id

*Matrix Multiplication with Expert Selection (Graph)***Description**

Indirect matrix multiplication for Mixture of Experts architectures. Selects expert weights based on indices and performs batched matmul.

**Usage**

```
ggml_mul_mat_id(ctx, as, b, ids)
```

**Arguments**

ctx	GGML context
as	Stacked expert weight matrices [n_embd, n_ff, nExperts]
b	Input tensor
ids	Expert selection indices tensor (I32)

**Value**

Output tensor after expert-selected matrix multiplication

**Examples**

```
## Not run:
ctx <- ggml_init(64 * 1024 * 1024)
# 4 experts, each with 8x16 weights (small for example)
experts <- ggml_new_tensor_3d(ctx, GGML_TYPE_F32, 8, 16, 4)
ggml_set_f32(experts, rnorm(8 * 16 * 4))
input <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 8, 2)
ggml_set_f32(input, rnorm(16))
# Select expert 0 for token 0, expert 2 for token 1
ids <- ggml_new_tensor_1d(ctx, GGML_TYPE_I32, 2)
ggml_set_i32(ids, c(0L, 2L))
output <- ggml_mul_mat_id(ctx, experts, input, ids)
graph <- ggml_build_forward_expand(ctx, output)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

*ggml\_nbytes*

*Get number of bytes*

**Description**

Get number of bytes

Get Number of Bytes

**Usage**

`ggml_nbytes(tensor)`

`ggml_nbytes(tensor)`

**Arguments**

<code>tensor</code>	Tensor
---------------------	--------

**Value**

Integer number of bytes

Integer number of bytes

## Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
tensor <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
ggml nbytes(tensor)  
ggml_free(ctx)  
  
## End(Not run)  
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
ggml nbytes(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_neg

*Negation (Graph)*

---

## Description

Creates a graph node for element-wise negation: -x

## Usage

```
ggml_neg(ctx, a)
```

## Arguments

ctx	GGML context
a	Input tensor

## Value

Tensor representing the negation operation

---

ggml\_nelements

*Get number of elements*

---

## Description

Get number of elements

Get Number of Elements

**Usage**

```
ggml_nelements(tensor)

ggml_nelements(tensor)
```

**Arguments**

<code>tensor</code>	Tensor
---------------------	--------

**Value**

Integer number of elements
Integer number of elements

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
tensor <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 20)
ggml_nelements(tensor)
ggml_free(ctx)

## End(Not run)
## Not run:
ctx <- ggml_init(1024 * 1024)
t <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 20)
ggml_nelements(t)
ggml_free(ctx)

## End(Not run)
```

**ggml\_nelements***Create Scalar F32 Tensor***Description**

Creates a 1-element tensor containing a single float value. Useful for scalar operations, learning rates, and other scalar floats.

**Usage**

```
ggml_new_f32(ctx, value)
```

**Arguments**

<code>ctx</code>	GGML context
<code>value</code>	Numeric value

**Value**

Tensor pointer (1-element F32 tensor)

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
scalar <- ggml_new_f32(ctx, 3.14)  
ggml_get_f32(scalar)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_new\_i32

*Create Scalar I32 Tensor*

---

**Description**

Creates a 1-element tensor containing a single integer value. Useful for indices, counters, and other scalar integer operations.

**Usage**

```
ggml_new_i32(ctx, value)
```

**Arguments**

ctx	GGML context
value	Integer value

**Value**

Tensor pointer (1-element I32 tensor)

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
scalar <- ggml_new_i32(ctx, 42)  
ggml_get_i32(scalar)  
ggml_free(ctx)  
  
## End(Not run)
```

`ggml_new_tensor`      *Create Tensor with Arbitrary Dimensions*

### Description

Generic tensor constructor for creating tensors with 1-4 dimensions. This is more flexible than the `ggml_new_tensor_Nd` functions.

### Usage

```
ggml_new_tensor(ctx, type = GGML_TYPE_F32, n_dims, ne)
```

### Arguments

<code>ctx</code>	GGML context
<code>type</code>	Data type (GGML_TYPE_F32, etc.)
<code>n_dims</code>	Number of dimensions (1-4)
<code>ne</code>	Numeric vector of dimension sizes

### Value

Tensor pointer

### Examples

```
## Not run:
ctx <- ggml_init(1024 * 1024)
t <- ggml_new_tensor(ctx, GGML_TYPE_F32, 3, c(10, 20, 30))
ggml_nelements(t)
ggml_free(ctx)

## End(Not run)
```

`ggml_new_tensor_1d`      *Create 1D tensor*

### Description

Create 1D tensor

Create 1D Tensor

### Usage

```
ggml_new_tensor_1d(ctx, type = GGML_TYPE_F32, ne0)
```

```
ggml_new_tensor_1d(ctx, type = GGML_TYPE_F32, ne0)
```

**Arguments**

ctx	GGML context
type	Data type
ne0	Size

**Value**

Tensor pointer  
Tensor pointer

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
tensor <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)
ggml_nelements(tensor)
ggml_free(ctx)

## End(Not run)
```

ggml\_new\_tensor\_2d      *Create 2D tensor*

**Description**

Create 2D tensor  
Create 2D Tensor

**Usage**

```
ggml_new_tensor_2d(ctx, type = GGML_TYPE_F32, ne0, ne1)

ggml_new_tensor_2d(ctx, type = GGML_TYPE_F32, ne0, ne1)
```

**Arguments**

ctx	GGML context
type	Data type
ne0	Rows
ne1	Columns

**Value**

Tensor pointer  
Tensor pointer

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
tensor <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 20)
ggml_nelements(tensor)
ggml_free(ctx)

## End(Not run)
```

**ggml\_new\_tensor\_3d**      *Create 3D Tensor*

**Description**

Create 3D Tensor

**Usage**

```
ggml_new_tensor_3d(ctx, type = GGML_TYPE_F32, ne0, ne1, ne2)
```

**Arguments**

ctx	GGML context
type	Data type (default GGML_TYPE_F32)
ne0	Size of dimension 0
ne1	Size of dimension 1
ne2	Size of dimension 2

**Value**

Tensor pointer

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
t <- ggml_new_tensor_3d(ctx, GGML_TYPE_F32, 10, 20, 30)
ggml_nelements(t)
ggml_free(ctx)

## End(Not run)
```

---

ggml\_new\_tensor\_4d      *Create 4D Tensor*

---

**Description**

Create 4D Tensor

**Usage**

```
ggml_new_tensor_4d(ctx, type = GGML_TYPE_F32, ne0, ne1, ne2, ne3)
```

**Arguments**

ctx	GGML context
type	Data type (default GGML_TYPE_F32)
ne0	Size of dimension 0
ne1	Size of dimension 1
ne2	Size of dimension 2
ne3	Size of dimension 3

**Value**

Tensor pointer

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_4d(ctx, GGML_TYPE_F32, 8, 8, 3, 2)  
ggml_nelements(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_norm      *Layer Normalization (Graph)*

---

**Description**

Creates a graph node for layer normalization.

**Usage**

```
ggml_norm(ctx, a, eps = 1e-05)
```

**Arguments**

ctx	GGML context
a	Input tensor
eps	Epsilon value for numerical stability (default: 1e-5)

**Value**

Tensor representing the layer normalization operation

**ggml\_norm\_inplace**      *Layer Normalization In-place (Graph)*

**Description**

Creates a graph node for in-place layer normalization. Returns a view of the input tensor.

**Usage**

```
ggml_norm_inplace(ctx, a, eps = 1e-05)
```

**Arguments**

ctx	GGML context
a	Input tensor (will be modified in-place)
eps	Epsilon value for numerical stability (default: 1e-5)

**Value**

View of input tensor with layer normalization applied

**ggml\_nrows**      *Get Number of Rows*

**Description**

Returns the number of rows in a tensor (product of all dimensions except ne[0]).

**Usage**

```
ggml_nrows(tensor)
```

**Arguments**

tensor	Tensor pointer
--------	----------------

**Value**

Number of rows

---

ggml_n_dims	<i>Get Number of Dimensions</i>
-------------	---------------------------------

---

### Description

Returns the number of dimensions of a tensor

### Usage

```
ggml_n_dims(tensor)
```

### Arguments

tensor	Tensor pointer
--------	----------------

### Value

Integer number of dimensions (1-4)

### Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 20)  
ggml_n_dims(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml_out_prod	<i>Outer Product (Graph)</i>
---------------	------------------------------

---

### Description

Computes the outer product of two vectors:  $C = a * b^T$  For vectors  $a[m]$  and  $b[n]$ , produces matrix  $C[m, n]$ .

### Usage

```
ggml_out_prod(ctx, a, b)
```

### Arguments

ctx	GGML context
a	First vector tensor
b	Second vector tensor

**Value**

Matrix tensor representing the outer product

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 3)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 4)
ggml_set_f32(a, c(1, 2, 3))
ggml_set_f32(b, c(1, 2, 3, 4))
c <- ggml_out_prod(ctx, a, b) # Result: 3x4 matrix
graph <- ggml_build_forward_expand(ctx, c)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

ggml\_pad

*Pad Tensor with Zeros (Graph)***Description**

Pads tensor dimensions with zeros on the right side. Useful for aligning tensor sizes in attention operations.

**Usage**

```
ggml_pad(ctx, a, p0 = 0L, p1 = 0L, p2 = 0L, p3 = 0L)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor to pad
<code>p0</code>	Padding for dimension 0 (default 0)
<code>p1</code>	Padding for dimension 1 (default 0)
<code>p2</code>	Padding for dimension 2 (default 0)
<code>p3</code>	Padding for dimension 3 (default 0)

**Value**

Padded tensor with shape [ne0+p0, ne1+p1, ne2+p2, ne3+p3]

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 5, 3)
ggml_set_f32(a, 1:15)
# Pad to 8x4
b <- ggml_pad(ctx, a, 3, 1) # Add 3 zeros to dim0, 1 to dim1
graph <- ggml_build_forward_expand(ctx, b)
ggml_graph_compute(ctx, graph)
# Result shape: [8, 4]
ggml_free(ctx)

## End(Not run)
```

ggml\_permute

*Permute Tensor Dimensions (Graph)*

## Description

Permutes the tensor dimensions according to specified axes. CRITICAL for attention mechanisms in transformers.

## Usage

```
ggml_permute(ctx, a, axis0, axis1, axis2, axis3)
```

## Arguments

ctx	GGML context
a	Input tensor
axis0	New position for axis 0
axis1	New position for axis 1
axis2	New position for axis 2
axis3	New position for axis 3

## Value

Permuted tensor

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
# Create 4D tensor: (2, 3, 4, 5)
t <- ggml_new_tensor_4d(ctx, GGML_TYPE_F32, 2, 3, 4, 5)
# Swap axes 0 and 1: result shape (3, 2, 4, 5)
t_perm <- ggml_permute(ctx, t, 1, 0, 2, 3)
```

```
ggml_free(ctx)
## End(Not run)
```

**ggml\_pool\_1d***ID Pooling (Graph)***Description**

Applies 1D pooling operation.

**Usage**

```
ggml_pool_1d(ctx, a, op, k0, s0 = k0, p0 = 0L)
```

GGML\_OP\_POOL\_MAX

GGML\_OP\_POOL\_AVG

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor
<code>op</code>	Pool operation: GGML_OP_POOL_MAX (0) or GGML_OP_POOL_AVG (1)
<code>k0</code>	Kernel size
<code>s0</code>	Stride (default = <code>k0</code> )
<code>p0</code>	Padding (default 0)

**Format**

An object of class `integer` of length 1.

An object of class `integer` of length 1.

**Value**

Pooled tensor

---

`ggml_pool_2d`*2D Pooling (Graph)*

---

**Description**

Applies 2D pooling operation.

**Usage**

```
ggml_pool_2d(ctx, a, op, k0, k1, s0 = k0, s1 = k1, p0 = 0, p1 = 0)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor
<code>op</code>	Pool operation: GGML_OP_POOL_MAX (0) or GGML_OP_POOL_AVG (1)
<code>k0</code>	Kernel size dimension 0
<code>k1</code>	Kernel size dimension 1
<code>s0</code>	Stride dimension 0 (default = k0)
<code>s1</code>	Stride dimension 1 (default = k1)
<code>p0</code>	Padding dimension 0 (default 0)
<code>p1</code>	Padding dimension 1 (default 0)

**Value**

Pooled tensor

---

`ggml_print_mem_status` *Print Context Memory Status*

---

**Description**

Helper to print memory usage information

**Usage**

```
ggml_print_mem_status(ctx)
```

**Arguments**

<code>ctx</code>	GGML context
------------------	--------------

**Value**

List with total, used, free memory (invisible)

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
ggml_print_mem_status(ctx)
ggml_free(ctx)

## End(Not run)
```

`ggml_print_objects`      *Print Objects in Context*

**Description**

Debug function to print all objects (tensors) in the context

**Usage**

```
ggml_print_objects(ctx)
```

**Arguments**

<code>ctx</code>	GGML context
------------------	--------------

**Value**

NULL (invisible)

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
t <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)
ggml_print_objects(ctx)
ggml_free(ctx)

## End(Not run)
```

---

**ggml\_quantize\_chunk**    *Quantize Data Chunk*

---

**Description**

Quantizes a chunk of floating-point data to a lower precision format.

**Usage**

```
ggml_quantize_chunk(type, src, nrows, n_per_row)
```

**Arguments**

type	Target GGML type (e.g., GGML_TYPE_Q4_0)
src	Source numeric vector (F32 data)
nrows	Number of rows
n_per_row	Number of elements per row

**Value**

Raw vector containing quantized data

**Examples**

```
## Not run:  
# Quantize 256 floats to Q8_0 (block size 32)  
data <- rnorm(256)  
quantized <- ggml_quantize_chunk(GGML_TYPE_Q8_0, data, 1, 256)  
ggml_quantize_free() # Clean up  
  
## End(Not run)
```

---

**ggml\_quantize\_free**    *Free Quantization Resources*

---

**Description**

Frees any memory allocated by quantization. Call at end of program to avoid memory leaks.

**Usage**

```
ggml_quantize_free()
```

**Value**

NULL invisibly

---

`ggml_quantize_init`      *Initialize Quantization Tables*

---

### Description

Initializes quantization tables for a given type. Called automatically by `ggml_quantize_chunk`, but can be called manually.

### Usage

```
ggml_quantize_init(type)
```

### Arguments

`type`            GGML type (e.g., `GGML_TYPE_Q4_0`)

### Value

NULL invisibly

---

`ggml_quantize_requires_imatrix`  
Check if Quantization Requires Importance Matrix

---

### Description

Some quantization types require an importance matrix for optimal quality.

### Usage

```
ggml_quantize_requires_imatrix(type)
```

### Arguments

`type`            GGML type

### Value

TRUE if importance matrix is required

---

**ggml\_reglu***ReLU Gated Linear Unit) (Graph)*

---

## Description

Creates a graph node for ReGLU operation. ReGLU uses ReLU as the activation function on the first half.

## Usage

```
ggml_reglu(ctx, a)
```

## Arguments

ctx	GGML context
a	Input tensor (first dimension must be even)

## Details

Formula: output = ReLU(x) \* gate

## Value

Tensor with half the first dimension of input

## Examples

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
a <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 8, 3)  
ggml_set_f32(a, rnorm(24))  
r <- ggml_reglu(ctx, a)  
graph <- ggml_build_forward_expand(ctx, r)  
ggml_graph_compute(ctx, graph)  
result <- ggml_get_f32(r) # Shape: 4x3  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_reglu\_split      *ReGLU Split (Graph)*

---

### Description

Creates a graph node for ReGLU with separate input and gate tensors.

### Usage

```
ggml_reglu_split(ctx, a, b)
```

### Arguments

ctx	GGML context
a	Input tensor (the values to be gated)
b	Gate tensor (same shape as a)

### Details

Formula: output = ReLU(a) \* b

### Value

Tensor with same shape as input tensors

---

ggml\_relu      *ReLU Activation (Graph)*

---

### Description

Creates a graph node for ReLU (Rectified Linear Unit) activation: max(0, x)

### Usage

```
ggml_relu(ctx, a)
```

### Arguments

ctx	GGML context
a	Input tensor

### Value

Tensor representing the ReLU operation

## Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
ggml_set_f32(a, c(-2, -1, 0, 1, 2))  
result <- ggml_relu(ctx, a)  
graph <- ggml_build_forward_expand(ctx, result)  
ggml_graph_compute(ctx, graph)  
ggml_get_f32(result)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_repeat

*Repeat (Graph)*

---

## Description

Creates a graph node that repeats tensor 'a' to match shape of tensor 'b'.

## Usage

```
ggml_repeat(ctx, a, b)
```

## Arguments

ctx	GGML context
a	Tensor to repeat
b	Target tensor (defines output shape)

## Value

Tensor with repeated values

---

ggml\_repeat\_back

*Repeat Backward (Graph)*

---

## Description

Backward pass for repeat operation - sums repetitions back to original shape. Used for gradient computation during training.

## Usage

```
ggml_repeat_back(ctx, a, b)
```

**Arguments**

ctx	GGML context
a	Input tensor (gradients from repeated tensor)
b	Target shape tensor (original tensor before repeat)

**Value**

Tensor with summed gradients matching shape of b

ggml_reset	<i>Reset GGML Context</i>
------------	---------------------------

**Description**

Clears all tensor allocations in the context memory pool. The context can be reused without recreating it. This is more efficient than free + init for temporary operations.

**Usage**

```
ggml_reset(ctx)
```

**Arguments**

ctx	GGML context pointer
-----	----------------------

**Value**

NULL (invisible)

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 100)
ggml_reset(ctx)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 200)
ggml_free(ctx)

## End(Not run)
```

---

**ggml\_reshape\_1d**      *Reshape to 1D (Graph)*

---

**Description**

Reshapes tensor to 1D with ne0 elements

**Usage**

```
ggml_reshape_1d(ctx, a, ne0)
```

**Arguments**

ctx	GGML context
a	Input tensor
ne0	Size of dimension 0

**Value**

Reshaped tensor

---

**ggml\_reshape\_2d**      *Reshape to 2D (Graph)*

---

**Description**

Reshapes tensor to 2D with shape (ne0, ne1)

**Usage**

```
ggml_reshape_2d(ctx, a, ne0, ne1)
```

**Arguments**

ctx	GGML context
a	Input tensor
ne0	Size of dimension 0
ne1	Size of dimension 1

**Value**

Reshaped tensor

---

`ggml_reshape_3d`      *Reshape to 3D (Graph)*

---

**Description**

Reshapes tensor to 3D with shape (ne0, ne1, ne2)

**Usage**

```
ggml_reshape_3d(ctx, a, ne0, ne1, ne2)
```

**Arguments**

ctx	GGML context
a	Input tensor
ne0	Size of dimension 0
ne1	Size of dimension 1
ne2	Size of dimension 2

**Value**

Reshaped tensor

---

`ggml_reshape_4d`      *Reshape to 4D (Graph)*

---

**Description**

Reshapes tensor to 4D with shape (ne0, ne1, ne2, ne3)

**Usage**

```
ggml_reshape_4d(ctx, a, ne0, ne1, ne2, ne3)
```

**Arguments**

ctx	GGML context
a	Input tensor
ne0	Size of dimension 0
ne1	Size of dimension 1
ne2	Size of dimension 2
ne3	Size of dimension 3

**Value**

Reshaped tensor

---

ggml_rms_norm	<i>RMS Normalization (Graph)</i>
---------------	----------------------------------

---

## Description

Creates a graph node for RMS (Root Mean Square) normalization. CRITICAL for LLaMA models.

## Usage

```
ggml_rms_norm(ctx, a, eps = 1e-05)
```

## Arguments

ctx	GGML context
a	Input tensor
eps	Epsilon value for numerical stability (default: 1e-5)

## Value

Tensor representing the RMS normalization operation

---

ggml_rms_norm_back	<i>RMS Norm Backward (Graph)</i>
--------------------	----------------------------------

---

## Description

Creates a graph node for backward pass of RMS normalization. Used in training for computing gradients.

## Usage

```
ggml_rms_norm_back(ctx, a, b, eps = 1e-05)
```

## Arguments

ctx	GGML context
a	Input tensor (x from forward pass)
b	Gradient tensor (dy)
eps	Epsilon for numerical stability (default 1e-5)

## Value

Tensor representing the gradient with respect to input

`ggml_rms_norm_inplace` *RMS Normalization In-place (Graph)*

## Description

Creates a graph node for in-place RMS normalization. Returns a view of the input tensor. CRITICAL for LLaMA models when memory efficiency is important.

## Usage

```
ggml_rms_norm_inplace(ctx, a, eps = 1e-05)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (will be modified in-place)
<code>eps</code>	Epsilon value for numerical stability (default: 1e-5)

## Value

View of input tensor with RMS normalization applied

`ggml_rope`

*Rotary Position Embedding (Graph)*

## Description

Creates a graph node for RoPE (Rotary Position Embedding). RoPE is the dominant position encoding method in modern LLMs like LLaMA, Mistral, and many others.

## Usage

```
ggml_rope(ctx, a, b, n_dims, mode = 0L)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor of shape [head_dim, n_head, seq_len, batch]
<code>b</code>	Position tensor (int32) of shape [seq_len] containing position indices
<code>n_dims</code>	Number of dimensions to apply rotation to (usually head_dim)
<code>mode</code>	RoPE mode: GGML_ROPE_TYPE_NORM (0), GGML_ROPE_TYPE_NEON (2), etc.

## Details

RoPE encodes position information by rotating pairs of dimensions in the embedding space. The rotation angle depends on position and dimension index.

Key benefits of RoPE: - Relative position information emerges naturally from rotation - Better extrapolation to longer sequences than absolute embeddings - No additional parameters needed

## Value

Tensor with same shape as input, with rotary embeddings applied

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
# Query tensor: head_dim=8, n_head=4, seq_len=16, batch=1
q <- ggml_new_tensor_4d(ctx, GGML_TYPE_F32, 8, 4, 16, 1)
ggml_set_f32(q, rnorm(8 * 4 * 16))
# Position indices
pos <- ggml_new_tensor_1d(ctx, GGML_TYPE_I32, 16)
ggml_set_i32(pos, 0:15)
# Apply RoPE
q_rope <- ggml_rope(ctx, q, pos, 8, GGML_ROPE_TYPE_NORM)
graph <- ggml_build_forward_expand(ctx, q_rope)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

ggml\_rope\_ext

*Extended RoPE with Frequency Scaling (Graph)*

## Description

Creates a graph node for extended RoPE with frequency scaling parameters. Supports context extension techniques like YaRN, Linear Scaling, etc.

## Usage

```
ggml_rope_ext(
  ctx,
  a,
  b,
  c = NULL,
  n_dims,
  mode = 0L,
  n_ctx_orig = 0L,
  freq_base = 10000,
  freq_scale = 1,
```

```

    ext_factor = 0,
    attn_factor = 1,
    beta_fast = 32,
    beta_slow = 1
)

```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor
<code>b</code>	Position tensor (int32)
<code>c</code>	Optional frequency factors tensor (NULL for default)
<code>n_dims</code>	Number of dimensions to apply rotation to
<code>mode</code>	RoPE mode
<code>n_ctx_orig</code>	Original context length the model was trained on
<code>freq_base</code>	Base frequency for RoPE (default 10000 for most models)
<code>freq_scale</code>	Frequency scale factor (1.0 = no scaling)
<code>ext_factor</code>	YaRN extension factor (0.0 to disable)
<code>attn_factor</code>	Attention scale factor (typically 1.0)
<code>beta_fast</code>	YaRN parameter for fast dimensions
<code>beta_slow</code>	YaRN parameter for slow dimensions

### Details

This extended version supports various context extension techniques:

- **Linear Scaling**: Set `freq_scale = original_ctx / new_ctx` - **YaRN**: Set `ext_factor > 0` with appropriate `beta_fast/beta_slow` - **NTK-aware**: Adjust `freq_base` for NTK-style scaling

Common `freq_base` values: - LLaMA 1/2: 10000 - LLaMA 3: 500000 - Mistral: 10000 - Phi-3: 10000

### Value

Tensor with extended RoPE applied

### Examples

```

## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
q <- ggml_new_tensor_4d(ctx, GGML_TYPE_F32, 64, 8, 32, 1)
ggml_set_f32(q, rnorm(64 * 8 * 32))
pos <- ggml_new_tensor_1d(ctx, GGML_TYPE_I32, 32)
ggml_set_i32(pos, 0:31)
# Standard RoPE with default freq_base
q_rope <- ggml_rope_ext(ctx, q, pos, NULL,
                         n_dims = 64, mode = 0L,
                         n_ctx_orig = 4096,

```

```

freq_base = 10000, freq_scale = 1.0,
ext_factor = 0.0, attn_factor = 1.0,
beta_fast = 32, beta_slow = 1)
graph <- ggml_build_forward_expand(ctx, q_rope)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)

```

**ggml\_rope\_ext\_back**      *RoPE Extended Backward (Graph)*

### Description

Backward pass for extended RoPE (Rotary Position Embedding). Used during training to compute gradients through RoPE.

### Usage

```

ggml_rope_ext_back(
    ctx,
    a,
    b,
    c = NULL,
    n_dims,
    mode = 0L,
    n_ctx_orig = 0L,
    freq_base = 10000,
    freq_scale = 1,
    ext_factor = 0,
    attn_factor = 1,
    beta_fast = 32,
    beta_slow = 1
)

```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Gradient tensor from upstream (gradients of ggml_rope_ext result)
<code>b</code>	Position tensor (same as forward pass)
<code>c</code>	Optional frequency factors tensor (NULL for default)
<code>n_dims</code>	Number of dimensions for rotation
<code>mode</code>	RoPE mode
<code>n_ctx_orig</code>	Original context length
<code>freq_base</code>	Base frequency

<code>freq_scale</code>	Frequency scale factor
<code>ext_factor</code>	Extension factor (YaRN)
<code>attn_factor</code>	Attention factor
<code>beta_fast</code>	YaRN fast beta
<code>beta_slow</code>	YaRN slow beta

**Value**

Gradient tensor for the input

`ggml_rope_inplace`      *Rotary Position Embedding In-place (Graph)*

**Description**

In-place version of `ggml_rope`. Returns a view of the input tensor.

**Usage**

```
ggml_rope_inplace(ctx, a, b, n_dims, mode = 0L)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (will be modified in-place)
<code>b</code>	Position tensor (int32)
<code>n_dims</code>	Number of dimensions to apply rotation to
<code>mode</code>	RoPE mode

**Value**

View of input tensor with RoPE applied

---

`ggml_round`*Round (Graph)*

---

**Description**

Creates a graph node for element-wise rounding: round(x)

**Usage**

```
ggml_round(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the round operation

---

---

`ggml_scale`*Scale (Graph)*

---

**Description**

Creates a graph node for scaling tensor by a scalar: x \* s

**Usage**

```
ggml_scale(ctx, a, s)
```

**Arguments**

ctx	GGML context
a	Input tensor
s	Scalar value to multiply by

**Value**

Tensor representing the scaled values

<code>ggml_set</code>	<i>Set Tensor Region (Graph)</i>
-----------------------	----------------------------------

### Description

Copies tensor b into tensor a at a specified offset. This allows writing to a portion of a tensor.

### Usage

```
ggml_set(ctx, a, b, nb1, nb2, nb3, offset)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Destination tensor
<code>b</code>	Source tensor (data to copy)
<code>nb1</code>	Stride for dimension 1 (in bytes)
<code>nb2</code>	Stride for dimension 2 (in bytes)
<code>nb3</code>	Stride for dimension 3 (in bytes)
<code>offset</code>	Byte offset in destination tensor

### Value

Tensor representing the set operation

<code>ggml_set_1d</code>	<i>Set 1D Tensor Region (Graph)</i>
--------------------------	-------------------------------------

### Description

Simplified 1D version of `ggml_set`. Copies tensor b into tensor a starting at offset.

### Usage

```
ggml_set_1d(ctx, a, b, offset)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Destination tensor
<code>b</code>	Source tensor
<code>offset</code>	Byte offset in destination tensor

### Value

Tensor representing the set operation

---

ggml_set_2d	<i>Set 2D Tensor Region (Graph)</i>
-------------	-------------------------------------

---

**Description**

Simplified 2D version of ggml\_set.

**Usage**

```
ggml_set_2d(ctx, a, b, nb1, offset)
```

**Arguments**

ctx	GGML context
a	Destination tensor
b	Source tensor
nb1	Stride for dimension 1 (in bytes)
offset	Byte offset in destination tensor

**Value**

Tensor representing the set operation

---

ggml_set_f32	<i>Set F32 data</i>
--------------	---------------------

---

**Description**

Set F32 data

Set F32 Data

**Usage**

```
ggml_set_f32(tensor, data)
```

```
ggml_set_f32(tensor, data)
```

**Arguments**

tensor	Tensor
data	Numeric vector

**Value**

NULL (invisible)  
NULL (invisible)

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
tensor <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(tensor, c(1, 2, 3, 4, 5))
ggml_get_f32(tensor)
ggml_free(ctx)

## End(Not run)
## Not run:
ctx <- ggml_init(1024 * 1024)
t <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(t, c(1, 2, 3, 4, 5))
ggml_get_f32(t)
ggml_free(ctx)

## End(Not run)
```

ggml\_set\_i32

*Set I32 Data***Description**

Sets integer data in an I32 tensor. Used for indices (ggml\_get\_rows) and position tensors (ggml\_rope).

**Usage**

```
ggml_set_i32(tensor, data)
```

**Arguments**

tensor	Tensor of type GGML_TYPE_I32
data	Integer vector

**Value**

NULL (invisible)

## Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
pos <- ggml_new_tensor_1d(ctx, GGML_TYPE_I32, 10)  
ggml_set_i32(pos, 0:9)  
ggml_get_i32(pos)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_set\_name

*Set Tensor Name*

---

## Description

Assigns a name to a tensor. Useful for debugging and graph visualization.

## Usage

```
ggml_set_name(tensor, name)
```

## Arguments

tensor	Tensor pointer
name	Character string name

## Value

The tensor (for chaining)

---

ggml\_set\_no\_alloc

*Set No Allocation Mode*

---

## Description

When enabled, tensor creation will not allocate memory for data. Useful for creating computation graphs without allocating storage.

## Usage

```
ggml_set_no_alloc(ctx, no_alloc)
```

## Arguments

ctx	GGML context
no_alloc	Logical, TRUE to disable allocation

**Value**

NULL (invisible)

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
ggml_set_no_alloc(ctx, TRUE)
ggml_get_no_alloc(ctx)
ggml_set_no_alloc(ctx, FALSE)
ggml_free(ctx)

## End(Not run)
```

**ggml\_set\_n\_threads**      *Set Number of Threads*

**Description**

Set the number of threads for GGML operations

**Usage**

`ggml_set_n_threads(n_threads)`

**Arguments**

`n_threads`      Number of threads to use

**Value**

Number of threads set

**Examples**

```
## Not run:
# Use 4 threads
ggml_set_n_threads(4)

# Use all available cores
ggml_set_n_threads(parallel::detectCores())

## End(Not run)
```

---

ggml\_set\_zero            *Set Tensor to Zero*

---

### Description

Sets all elements of a tensor to zero. This is more efficient than manually setting all elements.

### Usage

```
ggml_set_zero(tensor)
```

### Arguments

tensor            Tensor to zero out

### Value

NULL (invisible)

### Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
ggml_set_f32(t, 1:10)  
ggml_set_zero(t)  
ggml_get_f32(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_sgn            *Sign Function (Graph)*

---

### Description

Creates a graph node for element-wise sign function.  $\text{sgn}(x) = -1$  if  $x < 0$ ,  $0$  if  $x == 0$ ,  $1$  if  $x > 0$

### Usage

```
ggml_sgn(ctx, a)
```

### Arguments

ctx            GGML context  
a            Input tensor

**Value**

Tensor representing the sign operation

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-2, -0.5, 0, 0.5, 2))
r <- ggml_sgn(ctx, a)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r) # c(-1, -1, 0, 1, 1)
ggml_free(ctx)

## End(Not run)
```

**ggml\_sigmoid**

*Sigmoid Activation (Graph)*

**Description**

Creates a graph node for sigmoid activation:  $1 / (1 + \exp(-x))$

**Usage**

```
ggml_sigmoid(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the sigmoid operation

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-2, -1, 0, 1, 2))
result <- ggml_sigmoid(ctx, a)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

---

ggml_silu	<i>SiLU Activation (Graph)</i>
-----------	--------------------------------

---

## Description

Creates a graph node for SiLU (Sigmoid Linear Unit) activation, also known as Swish. CRITICAL for LLaMA models.

## Usage

```
ggml_silu(ctx, a)
```

## Arguments

ctx	GGML context
a	Input tensor

## Value

Tensor representing the SiLU operation

## Examples

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
ggml_set_f32(a, c(-2, -1, 0, 1, 2))  
result <- ggml_silu(ctx, a)  
graph <- ggml_build_forward_expand(ctx, result)  
ggml_graph_compute(ctx, graph)  
ggml_get_f32(result)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml_silu_back	<i>SiLU Backward (Graph)</i>
----------------	------------------------------

---

## Description

Computes the backward pass for SiLU (Swish) activation. Used during training for gradient computation.

## Usage

```
ggml_silu_back(ctx, a, b)
```

**Arguments**

ctx	GGML context
a	Forward input tensor
b	Gradient tensor from upstream

**Value**

Gradient tensor for the input

---

ggml_sin	<i>Sine (Graph)</i>
----------	---------------------

---

**Description**

Creates a graph node for element-wise sine:  $\sin(x)$

**Usage**

```
ggml_sin(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the sin operation

---

ggml_softplus	<i>Softplus Activation (Graph)</i>
---------------	------------------------------------

---

**Description**

Creates a graph node for Softplus activation.  $\text{Softplus}(x) = \log(1 + \exp(x))$ . A smooth approximation of ReLU.

**Usage**

```
ggml_softplus(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the Softplus operation

**Examples**

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)  
ggml_set_f32(a, c(-2, -1, 0, 1, 2))  
r <- ggml_softplus(ctx, a)  
graph <- ggml_build_forward_expand(ctx, r)  
ggml_graph_compute(ctx, graph)  
result <- ggml_get_f32(r)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_soft\_max

*Softmax (Graph)*

---

**Description**

Creates a graph node for softmax operation. CRITICAL for attention mechanisms.

**Usage**

```
ggml_soft_max(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the softmax operation

`ggml_soft_max_ext`*Extended Softmax with Masking and Scaling (Graph)*

## Description

Creates a graph node for fused softmax operation with optional masking and ALiBi (Attention with Linear Biases) support. Computes: softmax( $a * \text{scale} + \text{mask} * (\text{ALiBi slope})$ ) CRITICAL for efficient attention computation in transformers.

## Usage

```
ggml_soft_max_ext(ctx, a, mask = NULL, scale = 1, max_bias = 0)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (typically attention scores)
<code>mask</code>	Optional attention mask tensor (F16 or F32). NULL for no mask. Shape must be broadcastable to input tensor.
<code>scale</code>	Scaling factor, typically $1/\sqrt{\text{head\_dim}}$
<code>max_bias</code>	Maximum ALiBi bias (0.0 to disable ALiBi)

## Details

This extended softmax is commonly used in transformer attention: 1. Scale attention scores by  $1/\sqrt{d_k}$  for numerical stability 2. Apply attention mask (e.g., causal mask, padding mask) 3. Optionally apply ALiBi position bias 4. Compute softmax

All these operations are fused for efficiency.

## Value

Tensor representing the scaled and masked softmax

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
scores <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 10)
ggml_set_f32(scores, rnorm(100))
attn <- ggml_soft_max_ext(ctx, scores, NULL, 1.0, max_bias = 0.0)
graph <- ggml_build_forward_expand(ctx, attn)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

---

**ggml\_soft\_max\_ext\_back***Softmax Backward Extended (Graph)*

---

**Description**

Backward pass for extended softmax operation.

**Usage**

```
ggml_soft_max_ext_back(ctx, a, b, scale = 1, max_bias = 0)
```

**Arguments**

ctx	GGML context
a	Softmax output tensor (from forward pass)
b	Gradient tensor from upstream
scale	Scale factor (same as forward pass)
max_bias	Maximum ALiBi bias (same as forward pass)

**Value**

Gradient tensor for the input

---

**ggml\_soft\_max\_inplace** *Softmax In-place (Graph)*

---

**Description**

Creates a graph node for in-place softmax operation. Returns a view of the input tensor.

**Usage**

```
ggml_soft_max_inplace(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor (will be modified in-place)

**Value**

View of input tensor with softmax applied

`GGML_SORT_ORDER_ASC`     *Sort Order Constants*

### Description

Constants for specifying sort order in argsort operations.

### Usage

```
GGML_SORT_ORDER_ASC
GGML_SORT_ORDER_DESC
```

### Format

Integer constants  
An object of class `integer` of length 1.

### Examples

```
## Not run:
GGML_SORT_ORDER_ASC # 0 - Ascending order
GGML_SORT_ORDER_DESC # 1 - Descending order

## End(Not run)
```

`ggml_sqr`                    *Square (Graph)*

### Description

Creates a graph node for element-wise squaring:  $x^2$

### Usage

```
ggml_sqr(ctx, a)
```

### Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor

### Value

Tensor representing the square operation

---

**ggml\_sqrt***Square Root (Graph)*

---

**Description**

Creates a graph node for element-wise square root:  $\text{sqrt}(x)$

**Usage**

```
ggml_sqrt(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the  $\text{sqrt}$  operation

---

---

**ggml\_step***Step Function (Graph)*

---

**Description**

Creates a graph node for element-wise step function.  $\text{step}(x) = 0 \text{ if } x \leq 0, 1 \text{ if } x > 0$  Also known as the Heaviside step function.

**Usage**

```
ggml_step(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor representing the step operation

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-2, -0.5, 0, 0.5, 2))
r <- ggml_step(ctx, a)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r) # c(0, 0, 0, 1, 1)
ggml_free(ctx)

## End(Not run)
```

ggml\_sub

*Element-wise Subtraction (Graph)***Description**

Creates a graph node for element-wise subtraction.

**Usage**

```
ggml_sub(ctx, a, b)
```

**Arguments**

ctx	GGML context
a	First tensor
b	Second tensor (same shape as a)

**Value**

Tensor representing the subtraction operation ( $a - b$ )

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
b <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(5, 4, 3, 2, 1))
ggml_set_f32(b, c(1, 1, 1, 1, 1))
result <- ggml_sub(ctx, a, b)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

---

`ggml_sum`*Sum (Graph)*

---

**Description**

Creates a graph node that computes the sum of all elements.

**Usage**

```
ggml_sum(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Scalar tensor with the sum

---

`ggml_sum_rows`*Sum Rows (Graph)*

---

**Description**

Creates a graph node that computes the sum along rows.

**Usage**

```
ggml_sum_rows(ctx, a)
```

**Arguments**

ctx	GGML context
a	Input tensor

**Value**

Tensor with row sums

`ggml_swiglu`*SwiGLU (Swish/SiLU Gated Linear Unit) (Graph)*

## Description

Creates a graph node for SwiGLU operation. SwiGLU uses SiLU (Swish) as the activation function on the first half. CRITICAL for LLaMA, Mistral, and many modern LLMs.

## Usage

```
ggml_swiglu(ctx, a)
```

## Arguments

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (first dimension must be even)

## Details

Formula:  $\text{output} = \text{SiLU}(x) * \text{gate}$

## Value

Tensor with half the first dimension of input

## Examples

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
a <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 8, 3)
ggml_set_f32(a, rnorm(24))
r <- ggml_swiglu(ctx, a)
graph <- ggml_build_forward_expand(ctx, r)
ggml_graph_compute(ctx, graph)
result <- ggml_get_f32(r) # Shape: 4x3
ggml_free(ctx)

## End(Not run)
```

---

ggml\_swiglu\_split      *SwiGLU Split (Graph)*

---

### Description

Creates a graph node for SwiGLU with separate input and gate tensors.

### Usage

```
ggml_swiglu_split(ctx, a, b)
```

### Arguments

ctx	GGML context
a	Input tensor (the values to be gated)
b	Gate tensor (same shape as a)

### Details

Formula: output = SiLU(a) \* b

### Value

Tensor with same shape as input tensors

---

ggml\_tanh      *Tanh Activation (Graph)*

---

### Description

Creates a graph node for hyperbolic tangent activation.

### Usage

```
ggml_tanh(ctx, a)
```

### Arguments

ctx	GGML context
a	Input tensor

### Value

Tensor representing the tanh operation

**Examples**

```
## Not run:
ctx <- ggml_init(1024 * 1024)
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 5)
ggml_set_f32(a, c(-2, -1, 0, 1, 2))
result <- ggml_tanh(ctx, a)
graph <- ggml_build_forward_expand(ctx, result)
ggml_graph_compute(ctx, graph)
ggml_get_f32(result)
ggml_free(ctx)

## End(Not run)
```

**ggml\_tensor\_overhead**    *Get Tensor Overhead*

**Description**

Returns the memory overhead (metadata) for each tensor in bytes

**Usage**

```
ggml_tensor_overhead()
```

**Value**

Size in bytes

**Examples**

```
## Not run:
ggml_tensor_overhead()

## End(Not run)
```

**ggml\_tensor\_shape**    *Get Tensor Shape*

**Description**

Returns the shape of a tensor as a numeric vector of 4 elements (ne0, ne1, ne2, ne3)

**Usage**

```
ggml_tensor_shape(tensor)
```

**Arguments**

tensor      Tensor pointer

**Value**

Numeric vector of length 4 with dimensions

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 10, 20)  
ggml_tensor_shape(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_tensor\_type      *Get Tensor Type*

---

**Description**

Returns the data type of a tensor as an integer code

**Usage**

ggml\_tensor\_type(tensor)

**Arguments**

tensor      Tensor pointer

**Value**

Integer type code (0 = F32, 1 = F16, etc.)

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
t <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
ggml_tensor_type(t)  
ggml_free(ctx)  
  
## End(Not run)
```

---

`ggml_test`*Test GGML*

---

**Description**

Runs GGML library self-test and prints version info.

**Usage**

```
ggml_test()
```

**Value**

TRUE if test passed

**Examples**

```
## Not run:  
ggml_test()  
  
## End(Not run)
```

---

`ggml_time_init`*Initialize GGML Timer*

---

**Description**

Initializes the GGML timing system. Call this once at the beginning of the program before using `ggml_time_ms()` or `ggml_time_us()`.

**Usage**

```
ggml_time_init()
```

**Value**

NULL (invisible)

**Examples**

```
## Not run:  
ggml_time_init()  
start <- ggml_time_ms()  
Sys.sleep(0.01)  
elapsed <- ggml_time_ms() - start  
  
## End(Not run)
```

---

ggml\_time\_ms

*Get Time in Milliseconds*

---

## Description

Returns the current time in milliseconds since the timer was initialized.

## Usage

```
ggml_time_ms()
```

## Value

Numeric value representing milliseconds

## Examples

```
## Not run:  
ggml_time_init()  
start <- ggml_time_ms()  
Sys.sleep(0.01)  
elapsed <- ggml_time_ms() - start  
  
## End(Not run)
```

---

ggml\_time\_us

*Get Time in Microseconds*

---

## Description

Returns the current time in microseconds since the timer was initialized. More precise than ggml\_time\_ms() for micro-benchmarking.

## Usage

```
ggml_time_us()
```

## Value

Numeric value representing microseconds

**Examples**

```
## Not run:
ggml_time_init()
start <- ggml_time_us()
Sys.sleep(0.001)
elapsed <- ggml_time_us() - start

## End(Not run)
```

**ggml\_top\_k***Top-K Indices (Graph)***Description**

Returns the indices of top K elements per row. Useful for sampling strategies in language models (top-k sampling). Note: the resulting indices are in no particular order within top-k.

**Usage**

```
ggml_top_k(ctx, a, k)
```

**Arguments**

<code>ctx</code>	GGML context
<code>a</code>	Input tensor (F32)
<code>k</code>	Number of top elements to return per row

**Value**

Tensor containing I32 indices of top-k elements (not values)

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
# Logits from model output
logits <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 100)
ggml_set_f32(logits, rnorm(100))
# Get top 5 logits for sampling
top5 <- ggml_top_k(ctx, logits, 5)
graph <- ggml_build_forward_expand(ctx, top5)
ggml_graph_compute(ctx, graph)
ggml_free(ctx)

## End(Not run)
```

---

ggml_transpose	<i>Transpose (Graph)</i>
----------------	--------------------------

---

## Description

Creates a graph node for matrix transpose operation.

## Usage

```
ggml_transpose(ctx, a)
```

## Arguments

ctx	GGML context
a	Input tensor (2D matrix)

## Value

Tensor representing the transposed matrix

---

GGML_TYPE_F32	<i>GGML Data Types</i>
---------------	------------------------

---

## Description

Constants representing different data types supported by GGML.

## Usage

```
GGML_TYPE_F32  
GGML_TYPE_F16  
GGML_TYPE_Q4_0  
GGML_TYPE_Q4_1  
GGML_TYPE_Q8_0  
GGML_TYPE_I32
```

## Format

Integer constants  
 An object of class `integer` of length 1.  
 An object of class `integer` of length 1.

## Details

- `GGML_TYPE_F32`: 32-bit floating point (default)
- `GGML_TYPE_F16`: 16-bit floating point (half precision)
- `GGML_TYPE_Q4_0`: 4-bit quantization type 0
- `GGML_TYPE_Q4_1`: 4-bit quantization type 1
- `GGML_TYPE_Q8_0`: 8-bit quantization type 0
- `GGML_TYPE_I32`: 32-bit integer

## Examples

```
## Not run:  
GGML_TYPE_F32  
GGML_TYPE_F16  
GGML_TYPE_I32  
  
## End(Not run)
```

`ggml_type_size`      *Get Type Size in Bytes*

## Description

Returns the size in bytes for all elements in a block for a given type.

## Usage

`ggml_type_size(type)`

## Arguments

type	GGML type constant (e.g., <code>GGML_TYPE_F32</code> )
------	--

## Value

Size in bytes

---

ggml_upscale	<i>Upscale Tensor (Graph)</i>
--------------	-------------------------------

---

**Description**

Upscales tensor by multiplying ne0 and ne1 by scale factor. Supports different interpolation modes.

**Usage**

```
ggml_upscale(ctx, a, scale_factor, mode = 0L)

GGML_SCALE_MODE_NEAREST

GGML_SCALE_MODE_BILINEAR

GGML_SCALE_MODE_BICUBIC
```

**Arguments**

ctx	GGML context
a	Input tensor
scale_factor	Integer scale factor (e.g., 2 = double size)
mode	Scale mode: 0 = nearest, 1 = bilinear, 2 = bicubic

**Format**

An object of class `integer` of length 1.

An object of class `integer` of length 1.

An object of class `integer` of length 1.

**Value**

Upscaled tensor

**Examples**

```
## Not run:
ctx <- ggml_init(16 * 1024 * 1024)
img <- ggml_new_tensor_2d(ctx, GGML_TYPE_F32, 8, 8)
ggml_set_f32(img, rnorm(64))
upscaled <- ggml_upscale(ctx, img, 2, GGML_SCALE_MODE_NEAREST)
graph <- ggml_build_forward_expand(ctx, upscaled)
ggml_graph_compute(ctx, graph)
# Result is 16x16
ggml_free(ctx)

## End(Not run)
```

---

ggml\_used\_mem      *Get Used Memory*

---

**Description**

Returns the amount of memory currently used in the context

**Usage**

```
ggml_used_mem(ctx)
```

**Arguments**

ctx	GGML context
-----	--------------

**Value**

Used memory in bytes

**Examples**

```
## Not run:  
ctx <- ggml_init(1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 100)  
ggml_used_mem(ctx)  
ggml_free(ctx)  
  
## End(Not run)
```

---

ggml\_version      *Get GGML version*

---

**Description**

Get GGML version

**Usage**

```
ggml_version()
```

**Value**

Character string with GGML version

## Examples

```
## Not run:  
ggml_version()  
  
## End(Not run)
```

---

ggml\_view\_1d

*ID View with Byte Offset (Graph)*

---

## Description

Creates a 1D view of a tensor starting at a byte offset. The view shares memory with the source tensor.

## Usage

```
ggml_view_1d(ctx, a, ne0, offset = 0)
```

## Arguments

ctx	GGML context
a	Source tensor
ne0	Number of elements in the view
offset	Byte offset from the start of tensor data

## Value

View tensor

## Examples

```
## Not run:  
ctx <- ggml_init(16 * 1024 * 1024)  
a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 100)  
# View elements 10-19 (offset = 10 * 4 bytes = 40)  
v <- ggml_view_1d(ctx, a, 10, 40)  
ggml_free(ctx)  
  
## End(Not run)
```

---

`ggml_view_2d`*2D View with Byte Offset (Graph)*

---

## Description

Creates a 2D view of a tensor starting at a byte offset. The view shares memory with the source tensor.

## Usage

```
ggml_view_2d(ctx, a, ne0, ne1, nb1, offset = 0)
```

## Arguments

ctx	GGML context
a	Source tensor
ne0	Size of dimension 0
ne1	Size of dimension 1
nb1	Stride for dimension 1 (in bytes)
offset	Byte offset from the start of tensor data

## Value

View tensor

---

`ggml_view_3d`*3D View with Byte Offset (Graph)*

---

## Description

Creates a 3D view of a tensor starting at a byte offset. The view shares memory with the source tensor.

## Usage

```
ggml_view_3d(ctx, a, ne0, ne1, ne2, nb1, nb2, offset = 0)
```

**Arguments**

ctx	GGML context
a	Source tensor
ne0	Size of dimension 0
ne1	Size of dimension 1
ne2	Size of dimension 2
nb1	Stride for dimension 1 (in bytes)
nb2	Stride for dimension 2 (in bytes)
offset	Byte offset from the start of tensor data

**Value**

View tensor

---

ggml\_view\_4d

*4D View with Byte Offset (Graph)*

---

**Description**

Creates a 4D view of a tensor starting at a byte offset. The view shares memory with the source tensor. CRITICAL for KV-cache operations in transformers.

**Usage**

```
ggml_view_4d(ctx, a, ne0, ne1, ne2, ne3, nb1, nb2, nb3, offset = 0)
```

**Arguments**

ctx	GGML context
a	Source tensor
ne0	Size of dimension 0
ne1	Size of dimension 1
ne2	Size of dimension 2
ne3	Size of dimension 3
nb1	Stride for dimension 1 (in bytes)
nb2	Stride for dimension 2 (in bytes)
nb3	Stride for dimension 3 (in bytes)
offset	Byte offset from the start of tensor data

**Value**

View tensor

---

`ggml_view_tensor`      *View Tensor*

---

**Description**

Creates a view of the tensor (shares data, no copy)

**Usage**

```
ggml_view_tensor(ctx, src)
```

**Arguments**

<code>ctx</code>	GGML context
<code>src</code>	Source tensor

**Value**

View tensor (shares data with src)

---

`ggml_vulkan_available`    *Check if Vulkan support is available*

---

**Description**

Returns TRUE if the package was compiled with Vulkan support. To enable Vulkan, reinstall with:  
install.packages(..., configure.args = "-with-vulkan")

**Usage**

```
ggml_vulkan_available()
```

**Value**

Logical indicating if Vulkan is available

**Examples**

```
ggml_vulkan_available()
```

---

```
ggml_vulkan_backend_name  
Get Vulkan backend name
```

---

### Description

Returns the name of the Vulkan backend (includes device info).

### Usage

```
ggml_vulkan_backend_name(backend)
```

### Arguments

backend	Vulkan backend pointer
---------	------------------------

### Value

Character string with backend name

### Examples

```
## Not run:  
if (ggml_vulkan_available() && ggml_vulkan_device_count() > 0) {  
  backend <- ggml_vulkan_init(0)  
  print(ggml_vulkan_backend_name(backend))  
  ggml_vulkan_free(backend)  
}  
  
## End(Not run)
```

---

---

```
ggml_vulkan_device_count  
Get number of Vulkan devices
```

---

### Description

Returns the number of available Vulkan-capable GPU devices.

### Usage

```
ggml_vulkan_device_count()
```

### Value

Integer count of Vulkan devices (0 if Vulkan not available)

## Examples

```
## Not run:  
if (ggml_vulkan_available()) {  
    ggml_vulkan_device_count()  
}  
  
## End(Not run)
```

---

`ggml_vulkan_device_description`  
*Get Vulkan device description*

---

## Description

Returns a human-readable description of the specified Vulkan device.

## Usage

```
ggml_vulkan_device_description(device = 0L)
```

## Arguments

device	Device index (0-based)
--------	------------------------

## Value

Character string with device description

## Examples

```
## Not run:  
if (ggml_vulkan_available() && ggml_vulkan_device_count() > 0) {  
    ggml_vulkan_device_description(0)  
}  
  
## End(Not run)
```

---

```
ggml_vulkan_device_memory  
Get Vulkan device memory
```

---

### Description

Returns free and total memory for the specified Vulkan device.

### Usage

```
ggml_vulkan_device_memory(device = 0L)
```

### Arguments

device      Device index (0-based)

### Value

Named list with 'free' and 'total' memory in bytes

### Examples

```
## Not run:  
if (ggml_vulkan_available() && ggml_vulkan_device_count() > 0) {  
  mem <- ggml_vulkan_device_memory(0)  
  cat("Free:", mem$free / 1e9, "GB\n")  
  cat("Total:", mem$total / 1e9, "GB\n")  
}  
  
## End(Not run)
```

---

```
ggml_vulkan_free      Free Vulkan backend
```

---

### Description

Releases resources associated with the Vulkan backend.

### Usage

```
ggml_vulkan_free(backend)
```

### Arguments

backend      Vulkan backend pointer from ggml\_vulkan\_init()

**Value**

NULL (invisible)

**Examples**

```
## Not run:
if (ggml_vulkan_available() && ggml_vulkan_device_count() > 0) {
  backend <- ggml_vulkan_init(0)
  ggml_vulkan_free(backend)
}

## End(Not run)
```

**ggml\_vulkan\_init**      *Initialize Vulkan backend*

**Description**

Creates a Vulkan backend for the specified device. The backend must be freed with `ggml_vulkan_free()` when done.

**Usage**

```
ggml_vulkan_init(device = 0L)
```

**Arguments**

device	Device index (0-based, default 0)
--------	-----------------------------------

**Value**

Vulkan backend pointer

**Examples**

```
## Not run:
if (ggml_vulkan_available() && ggml_vulkan_device_count() > 0) {
  backend <- ggml_vulkan_init(0)
  print(ggml_vulkan_backend_name(backend))
  ggml_vulkan_free(backend)
}

## End(Not run)
```

---

```
ggml_vulkan_is_backend
```

*Check if backend is Vulkan*

---

## Description

Returns TRUE if the given backend is a Vulkan backend.

## Usage

```
ggml_vulkan_is_backend(backend)
```

## Arguments

backend	Backend pointer
---------	-----------------

## Value

Logical indicating if backend is Vulkan

## Examples

```
## Not run:  
if (ggml_vulkan_available() && ggml_vulkan_device_count() > 0) {  
    vk_backend <- ggml_vulkan_init(0)  
    cpu_backend <- ggml_backend_cpu_init()  
  
    ggml_vulkan_is_backend(vk_backend) # TRUE  
    ggml_vulkan_is_backend(cpu_backend) # FALSE  
  
    ggml_vulkan_free(vk_backend)  
    ggml_backend_free(cpu_backend)  
}  
  
## End(Not run)
```

---

```
ggml_vulkan_list_devices
```

*List all Vulkan devices*

---

## Description

Returns detailed information about all available Vulkan devices.

## Usage

```
ggml_vulkan_list_devices()
```

**Value**

List of device information (index, name, memory)

**Examples**

```
## Not run:
if (ggml_vulkan_available() && ggml_vulkan_device_count() > 0) {
  devices <- ggml_vulkan_list_devices()
  print(devices)
}

## End(Not run)
```

**ggml\_vulkan\_status**      *Print Vulkan status*

**Description**

Prints information about Vulkan availability and devices.

**Usage**

```
ggml_vulkan_status()
```

**Value**

NULL (invisible), prints status to console

**Examples**

```
ggml_vulkan_status()
```

**ggml\_with\_temp\_ctx**      *Execute with Temporary Context*

**Description**

Creates a temporary context, executes code, and frees it automatically. Useful when you need to create large temporary tensors.

**Usage**

```
ggml_with_temp_ctx(mem_size, expr)
```

**Arguments**

mem_size	Context memory size in bytes
expr	Expression to evaluate with the temporary context

**Value**

Result of the expression

**Examples**

```
## Not run:  
# Create tensors in temporary context  
result <- ggml_with_temp_ctx(1024 * 1024, {  
  a <- ggml_new_tensor_1d(ctx, GGML_TYPE_F32, 10)  
  ggml_set_f32(a, 1:10)  
  ggml_get_f32(a)  
})  
  
## End(Not run)
```

**Description**

Constants for RoPE (Rotary Position Embedding) modes.

**Usage**

```
GGML_ROPE_TYPE_NORM  
GGML_ROPE_TYPE_NEOX  
GGML_ROPE_TYPE_MROPE  
GGML_ROPE_TYPE_VISION
```

**Format**

- An object of class `integer` of length 1.
- An object of class `integer` of length 1.
- An object of class `integer` of length 1.
- An object of class `integer` of length 1.

**Details**

- GGML\_ROPE\_TYPE\_NORM (0): Standard RoPE as in original paper - GGML\_ROPE\_TYPE\_NEOX (2): GPT-NeoX style RoPE (interleaved differently) - GGML\_ROPE\_TYPE\_MROPE (8): Multi-RoPE for models like Qwen2-VL - GGML\_ROPE\_TYPE\_VISION (24): Vision model RoPE

# Index

- \* **datasets**
  - GGML\_GLU\_OP\_REGLU, 59
  - ggml\_pool\_1d, 90
  - GGML\_SORT\_ORDER\_ASC, 120
  - GGML\_TYPE\_F32, 131
  - ggml\_upscale, 133
  - rope\_types, 145
- \* **internal**
  - ggmlR-package, 6
  - ggml\_abs, 7
  - ggml\_add, 7
  - ggml\_add1, 8
  - ggml\_are\_same\_shape, 9
  - ggml\_argmax, 9
  - ggml\_argsort, 10
  - ggml\_backend\_alloc\_ctx\_tensors, 11
  - ggml\_backend\_buffer\_free, 11
  - ggml\_backend\_buffer\_get\_size, 12
  - ggml\_backend\_buffer\_name, 12
  - ggml\_backend\_cpu\_init, 13
  - ggml\_backend\_cpu\_set\_n\_threads, 13
  - ggml\_backend\_free, 14
  - ggml\_backend\_graph\_compute, 14
  - ggml\_backend\_name, 15
  - ggml\_backend\_sched\_alloc\_graph, 15
  - ggml\_backend\_sched\_free, 16
  - ggml\_backend\_sched\_get\_backend, 16
  - ggml\_backend\_sched\_get\_n\_backends, 17
  - ggml\_backend\_sched\_get\_n\_copies, 17
  - ggml\_backend\_sched\_get\_n\_splits, 18
  - ggml\_backend\_sched\_get\_tensor\_backend, 18
  - ggml\_backend\_sched\_graph\_compute, 19
  - ggml\_backend\_sched\_graph\_compute\_async, 20
  - ggml\_backend\_sched\_new, 20
  - ggml\_backend\_sched\_reserve, 21
  - ggml\_backend\_sched\_reset, 22
- ggml\_backend\_sched\_set\_tensor\_backend, 23
- ggml\_backend\_sched\_synchronize, 23
- ggml\_backend\_tensor\_get\_data, 24
- ggml\_backend\_tensor\_set\_data, 24
- ggml\_build\_forward\_expand, 25
- ggml\_ceil, 26
- ggml\_clamp, 26
- ggml\_concat, 27
- ggml\_cont, 28
- ggml\_conv\_1d, 28
- ggml\_conv\_2d, 29
- ggml\_conv\_transpose\_1d, 29
- ggml\_cos, 30
- ggml\_cpu\_add, 31
- ggml\_cpu\_mul, 31
- ggml\_cpy, 32
- ggml\_cycles, 33
- ggml\_cycles\_per\_ms, 34
- ggml\_diag, 34
- ggml\_diag\_mask\_inf, 35
- ggml\_diag\_mask\_inf\_inplace, 36
- ggml\_diag\_mask\_zero, 37
- ggml\_div, 37
- ggml\_dup, 38
- ggml\_dup\_tensor, 39
- ggml\_element\_size, 39
- ggml\_elu, 40
- ggml\_estimate\_memory, 41
- ggml\_exp, 41
- ggml\_flash\_attn\_back, 42
- ggml\_flash\_attn\_ext, 42
- ggml\_floor, 44
- ggml\_free, 44
- ggml\_gallocr\_alloc\_graph, 45
- ggml\_gallocr\_free, 46
- ggml\_gallocr\_get\_buffer\_size, 46
- ggml\_gallocr\_new, 47
- ggml\_gallocr\_reserve, 47

ggml\_geglu, 48  
 ggml\_geglu\_quick, 49  
 ggml\_geglu\_split, 49  
 ggml\_gelu, 50  
 ggml\_gelu\_erf, 50  
 ggml\_gelu\_quick, 51  
 ggml\_get\_f32, 52  
 ggml\_get\_i32, 53  
 ggml\_get\_max\_tensor\_size, 53  
 ggml\_get\_mem\_size, 54  
 ggml\_get\_n\_threads, 56  
 ggml\_get\_name, 55  
 ggml\_get\_no\_alloc, 55  
 ggml\_get\_rows, 56  
 ggml\_get\_rows\_back, 57  
 ggml\_glu, 58  
 GGML\_GLU\_OP\_GEGLU (GGML\_GLU\_OP\_REGLU), 59  
 GGML\_GLU\_OP\_GEGLU\_ERF  
     (GGML\_GLU\_OP\_REGLU), 59  
 GGML\_GLU\_OP\_GEGLU\_QUICK  
     (GGML\_GLU\_OP\_REGLU), 59  
 GGML\_GLU\_OP\_REGLU, 59  
 GGML\_GLU\_OP\_SWIGLU (GGML\_GLU\_OP\_REGLU), 59  
 GGML\_GLU\_OP\_SWIGLU\_OAI  
     (GGML\_GLU\_OP\_REGLU), 59  
 ggml\_glu\_split, 60  
 ggml\_graph\_compute, 60  
 ggml\_graph\_compute\_with\_ctx, 61  
 ggml\_graph\_dump\_dot, 62  
 ggml\_graph\_get\_tensor, 63  
 ggml\_graph\_n\_nodes, 64  
 ggml\_graph\_node, 63  
 ggml\_graph\_overhead, 64  
 ggml\_graph\_print, 65  
 ggml\_graph\_reset, 65  
 ggml\_group\_norm, 66  
 ggml\_group\_norm\_inplace, 66  
 ggml\_hardsigmoid, 67  
 ggml\_hardswish, 67  
 ggml\_im2col, 68  
 ggml\_init, 69  
 ggml\_init\_auto, 70  
 ggml\_is\_available, 70  
 ggml\_is\_contiguous, 71  
 ggml\_is\_permuted, 71  
 ggml\_is\_transposed, 72  
 ggml\_l2\_norm, 73  
 ggml\_l2\_norm\_inplace, 73  
 ggml\_leaky\_relu, 74  
 ggml\_log, 74  
 ggml\_mean, 75  
 ggml\_mul, 75  
 ggml\_mul\_mat, 76  
 ggml\_mul\_mat\_id, 77  
 ggml\_n\_dims, 87  
 ggml nbytes, 78  
 ggml\_neg, 79  
 ggml\_nelements, 79  
 ggml\_new\_f32, 80  
 ggml\_new\_i32, 81  
 ggml\_new\_tensor, 82  
 ggml\_new\_tensor\_1d, 82  
 ggml\_new\_tensor\_2d, 83  
 ggml\_new\_tensor\_3d, 84  
 ggml\_new\_tensor\_4d, 85  
 ggml\_norm, 85  
 ggml\_norm\_inplace, 86  
 ggml\_nrows, 86  
 GGML\_OP\_POOL\_AVG (ggml\_pool\_1d), 90  
 GGML\_OP\_POOL\_MAX (ggml\_pool\_1d), 90  
 ggml\_out\_prod, 87  
 ggml\_pad, 88  
 ggml\_permute, 89  
 ggml\_pool\_1d, 90  
 ggml\_pool\_2d, 91  
 ggml\_print\_mem\_status, 91  
 ggml\_print\_objects, 92  
 ggml\_quantize\_chunk, 93  
 ggml\_quantize\_free, 93  
 ggml\_quantize\_init, 94  
 ggml\_quantize\_requires\_imatrix, 94  
 ggml\_reglu, 95  
 ggml\_reglu\_split, 96  
 ggml\_relu, 96  
 ggml\_repeat, 97  
 ggml\_repeat\_back, 97  
 ggml\_reset, 98  
 ggml\_reshape\_1d, 99  
 ggml\_reshape\_2d, 99  
 ggml\_reshape\_3d, 100  
 ggml\_reshape\_4d, 100  
 ggml\_rms\_norm, 101  
 ggml\_rms\_norm\_back, 101  
 ggml\_rms\_norm\_inplace, 102

ggml\_rope, 102  
ggml\_rope\_ext, 103  
ggml\_rope\_ext\_back, 105  
ggml\_rope\_inplace, 106  
GGML\_ROPE\_TYPE\_MROPE (rope\_types), 145  
GGML\_ROPE\_TYPE\_NEOX (rope\_types), 145  
GGML\_ROPE\_TYPE\_NORM (rope\_types), 145  
GGML\_ROPE\_TYPE\_VISION (rope\_types), 145  
ggml\_round, 107  
ggml\_scale, 107  
GGML\_SCALE\_MODE\_BICUBIC (ggml\_upscale),  
    133  
GGML\_SCALE\_MODE\_BILINEAR  
    (ggml\_upscale), 133  
GGML\_SCALE\_MODE\_NEAREST (ggml\_upscale),  
    133  
ggml\_set, 108  
ggml\_set\_1d, 108  
ggml\_set\_2d, 109  
ggml\_set\_f32, 109  
ggml\_set\_i32, 110  
ggml\_set\_n\_threads, 112  
ggml\_set\_name, 111  
ggml\_set\_no\_alloc, 111  
ggml\_set\_zero, 113  
ggml\_sgn, 113  
ggml\_sigmoid, 114  
ggml\_silu, 115  
ggml\_silu\_back, 115  
ggml\_sin, 116  
ggml\_soft\_max, 117  
ggml\_soft\_max\_ext, 118  
ggml\_soft\_max\_ext\_back, 119  
ggml\_soft\_max\_inplace, 119  
ggml\_softplus, 116  
GGML\_SORT\_ORDER\_ASC, 120  
GGML\_SORT\_ORDER\_DESC  
    (GGML\_SORT\_ORDER\_ASC), 120  
ggml\_sqrt, 120  
ggml\_sqrt, 121  
ggml\_step, 121  
ggml\_sub, 122  
ggml\_sum, 123  
ggml\_sum\_rows, 123  
ggml\_swiglu, 124  
ggml\_swiglu\_split, 125  
ggml\_tanh, 125  
ggml\_tensor\_overhead, 126  
ggml\_tensor\_shape, 126  
ggml\_tensor\_type, 127  
ggml\_test, 128  
ggml\_time\_init, 128  
ggml\_time\_ms, 129  
ggml\_time\_us, 129  
ggml\_top\_k, 130  
ggml\_transpose, 131  
GGML\_TYPE\_F16 (GGML\_TYPE\_F32), 131  
GGML\_TYPE\_F32, 131  
GGML\_TYPE\_I32 (GGML\_TYPE\_F32), 131  
GGML\_TYPE\_Q4\_0 (GGML\_TYPE\_F32), 131  
GGML\_TYPE\_Q4\_1 (GGML\_TYPE\_F32), 131  
GGML\_TYPE\_Q8\_0 (GGML\_TYPE\_F32), 131  
ggml\_type\_size, 132  
ggml\_upscale, 133  
ggml\_used\_mem, 134  
ggml\_version, 134  
ggml\_view\_1d, 135  
ggml\_view\_2d, 136  
ggml\_view\_3d, 136  
ggml\_view\_4d, 137  
ggml\_view\_tensor, 138  
ggml\_vulkan\_available, 138  
ggml\_vulkan\_backend\_name, 139  
ggml\_vulkan\_device\_count, 139  
ggml\_vulkan\_device\_description, 140  
ggml\_vulkan\_device\_memory, 141  
ggml\_vulkan\_free, 141  
ggml\_vulkan\_init, 142  
ggml\_vulkan\_is\_backend, 143  
ggml\_vulkan\_list\_devices, 143  
ggml\_vulkan\_status, 144  
ggml\_with\_temp\_ctx, 144  
ggmlR (ggmlR-package), 6  
ggmlR-package, 6  
rope\_types, 145